Text to Matrix Generator*
User's Guide

Dimitrios Zeimpekis† Efstratios Gallopoulos‡

Department of Computer Engineering and Informatics,
University of Patras, Greece

April 2007

*Work conducted in the context of and supported in part by a University of Patras KARATHEODORI grant.

†e-mail: dsz@hpclab.ceid.upatras.gr, supported in part by a Bodossaki Foundation graduate fellowship.

‡e-mail: stratis@ceid.upatras.gr.

Contents

1 Introduction **1**

2 Istallation Instuctions **1**

3 Graphical User Interfaces **3**
3.1 Indexing module (`tmg_gui`) 3
3.2 Dimensionality Reduction module (`dr_gui`) 6
3.3 Retrieval module (`retrieval_gui`) 9
3.4 Clustering module (`clustering_gui`) 11
3.5 Classification module (`classification_gui`) 13

A Appendix: Demonstration of Use **19**
A.1 Indexing module (`tmg_gui`) 19
A.2 Dimensionality Reduction module (`dr_gui`) 26
A.3 Retrieval module (`retrieval_gui`) 30
A.4 Clustering module (`clustering_gui`) 34
A.5 Classification module (`classification_gui`) 39

B Appendix: Function Reference **43**
`about_tmg_gui` 43
`block_diagonalize` 44
`classification_gui` 45
`clsi` 46
`clustering_gui` 47
`cm` 48
`col_normalization` 49
`col_rearrange` 50
`compute_scat` 51
`create_kmeans_response` 52
`create_pddp_response` 53
`create_retrieval_response` 54
`diff_vector` 55
`dr_gui` 56
`ekmeans` 57
`entropy` 58
`get_node_scat` 59
`gui` 60
`init_tmg` 61
`knn_multi` 62
`knn_single` 63
`ks_selection` 64
`ks_selection1` 65
`llsf_multi` 66
`llsf_single` 67

lsa 68
make_clusters_multi 69
make_clusters_single 70
make_labels 71
make_val_inds 72
merge_dictionary 73
merge_tdms 74
myperms 75
open_file 76
opt_2means 77
pca 78
pca_mat 79
pca_mat_afun 80
pca_propack 81
pca_propack_Atransfunc 82
pca_propack_afun 83
pddp 84
pddp_2means 85
pddp_optcut 86
pddp_optcut_2means 87
pddp_optcutpd 88
ps_pdf2ascii 89
retrieval_gui 90
rocchio_multi 91
rocchio_single 92
scut_knn 93
scut_llsf 94
scut_rocchio 95
sdd_tmg 96
skmeans 97
stemmer 98
strip_html 99
svd_tmg 100
tdm_downdate 101
tdm_update 102
tmg 104
tmg_gui 106
tmg_query 107
tmg_save_results 109
tmg_template 110
two_means_1d 111
unique_elements 112
unique_words 113
vsm 114

List of Figures

1 Structure and dependencies of GUI modules of TMG. 1
2 Structure of TMG root directory. 2
3 The `tmg_gui` GUI. 3
4 The `dr_gui` GUI. 6
5 The `retrieval_gui` GUI. 9
6 The `clustering_gui` GUI. 11
7 The `classification_gui` GUI. 13
8 Starting window of `tmg_gui`. 20
9 Next view of `tmg_gui` according to the user selection. 21
10 The `open_file` window. 22
11 The output '.mat' files of `tmg_gui`. 23
12 The MySQL view uppon `tmg` execution. 24
13 The GUIs' general help tab. 25
14 Starting window of `dr_gui`. 27
15 Next view of `dr_gui` according to the user selection. 28
16 The output '.mat' files of `dr_gui`. 29
17 Starting window of `retrieval_gui`. 31
18 Next view of `retrieval_gui` according to the user selection. . . . 32
19 The output of `retrieval_gui`. 33
20 Starting window of `clustering_gui`. 35
21 Next view of `clustering_gui` according to the user selection. . . 36
22 The output '.mat' files of `clustering_gui`. 37
23 The output of `clustering_gui` for PDDP. 38
24 Starting window of `classification_gui`. 40
25 Next view of `classification_gui` according to the user selection. 41

List of Tables

1 Description of use of `tmg_gui` components. 5
2 Description of use of `dr_gui` components. 8
3 Description of use of `retrieval_gui` components. 10
4 Description of use of `clustering_gui` components. 12
5 Description of use of `classification_gui` components. 15

1 Introduction

Text to Matrix Generator (TMG) is a MATLAB Toolbox that can be used for various Data Mining (DM) and Information Retrieval (IR) tasks. TMG uses the sparse matrix infrastracture of MATLAB that is especially suited for TM applications where data are extremely sparse. Initially built as a preprocessing tool, TMG offers now a wide range of DM tools. In particular, TMG is composed of five Graphical User Interface (GUI) modules, presented in Figure 1 (arrows show modules' dependencies).

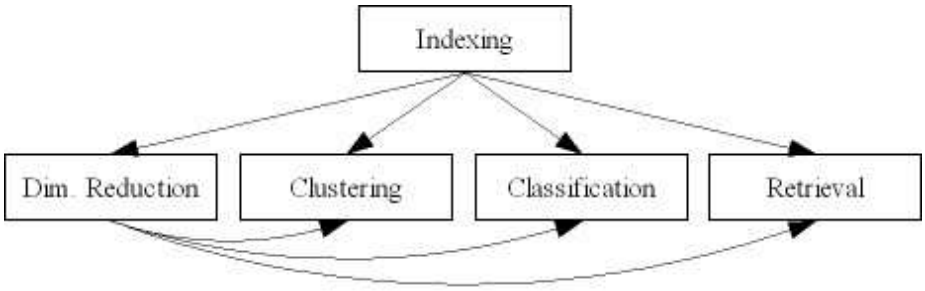

Figure 1: Structure and dependencies of GUI modules of TMG.

In the sequel, we first discuss the installation procedure of TMG and then describe in some detail the GUI's usage. In Appendix A we give a demonstation of use for all the TMG components, while Appendix B supplies a function reference.

2 Istallation Instuctions

Installation of TMG is straightforward by means of the `init_tmg` script. In particular, the user has to perform the following steps:

- For MySQL functionality, install MySQL and Java Connector.
- Download TMG by filling the form from: http://scgroup.hpclab.ceid.upatras.gr/scgroup/Projects/TMG/tmg_request.php[1]
- Unzip TMG_X.XRX.zip and start MATLAB. Figure 2 depicts the directory structure of the TMG root directory.
- Change path to the TMG root directory.
- Run init_tmg. Give the MySQL login and password as well as the root directory of the MySQL Java Connector. The installation script creates all necessary information (including MySQL database TMG) and adds to the MATLAB path all necessary directories.
- Run gui. Alternatively, use the command line interface, type `help tmg`.

[1] We are currently on the development of a wiki page.

TMG requires the MySQL[2], PROPACK[3], SDDPACK[4] and SPQR[5] third party software packages. PROPACK, SDDPACK and SPQR packages are included into TMG, while the user has to download MySQL. However, we note that MySQL related software is necessary only if the user intends to use the database support implemented into TMG. Ordinary TMG will run without any problem on a Matlab 7.0 environment without any other special software.

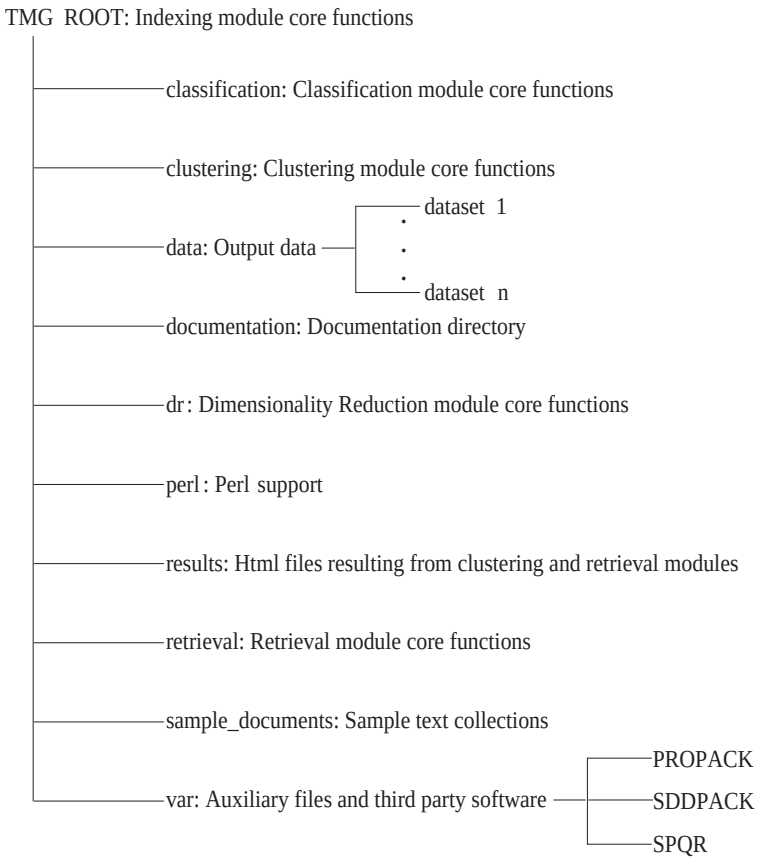

Figure 2: Structure of TMG root directory.

[2]http://www.mysql.com/, http://dev.mysql.com/downloads/connector/j/5.0.html
[3]http://soi.stanford.edu/rmunk/PROPACK/index.html
[4]http://www.cs.umd.edu/oleary/SDDPACK/README.html
[5]http://portal.acm.org/citation.cfm?id=1067972

3 Graphical User Interfaces

3.1 Indexing module (`tmg_gui`)

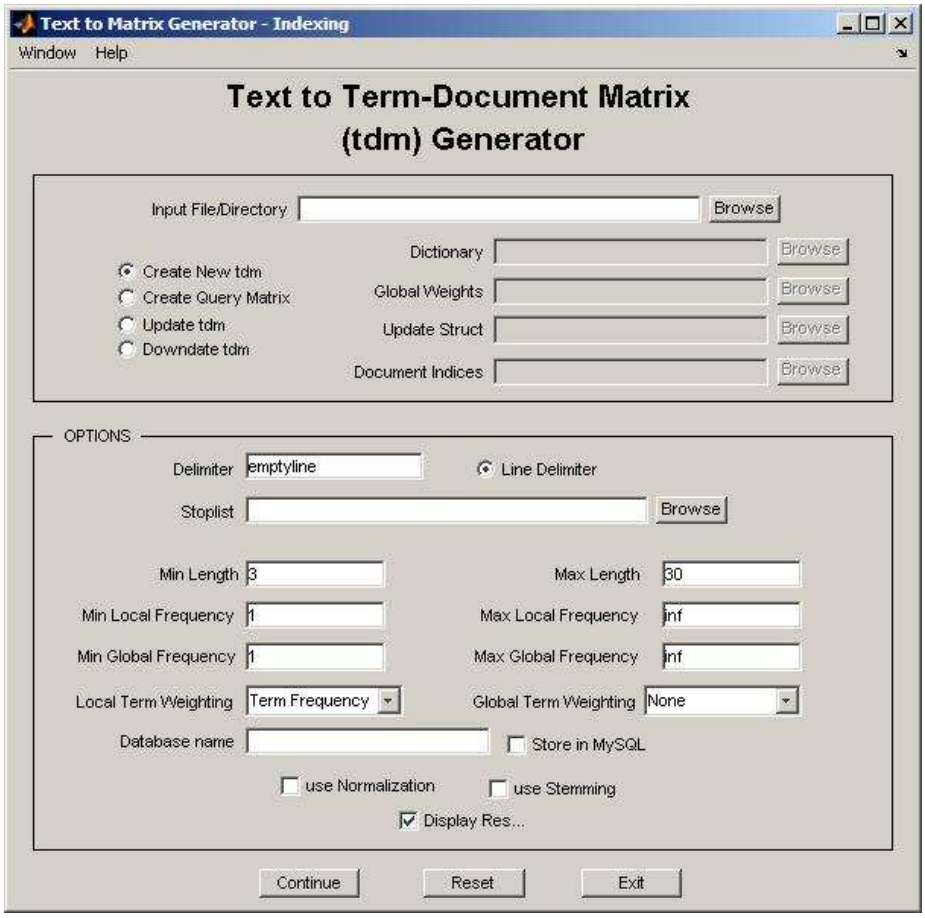

Figure 3: The `tmg_gui` GUI.

TMG can be used for the construction of new and the update of existing term-document matrices (tdms) from text collections, in the form of MATLAB sparse arrays. To this end, TMG implements various steps such as:

- Removal of stopwords.
- Apply stemming (currently Porter stemming algorithm [8]).
- Remove of short/long terms.
- Remove of frequent/infrequent terms (locally or globally).

- Term weighting and normalization.
- Html filtering, processing of Postscript and PDF.
- Store in MySQL (optionally).

The resulting tdms can be stored as "mat" files, while text can also be stored in MySQL for further procesing. TMG can also update existing tdms by efficient incremental updating or downdating operations. Finally, TMG can also construct query vectors using the existing dictionary that can be used from the retrieval and classification modules.

The indexing GUI module is depicted in Figure 3 while Table 1 describes in detail all the `tmg_gui` fields.

Field Name	Default	Description
Input File/Directory	-	Files to be parsed with resulting documents separated by "Delimiter". Alternatively, each file in the input directory contains a single document.
Create New tdm	•	Checked if new tdm is to be created (default checked).
Create Query Matrix	-	Checked if new query matrix is to be created (default checked).
Update tdm	-	Checked if an existing tdm is to be updated with new documents. Alternatively, ckecked if an existing tdm is to be updated using different options (change update_struct).
Downdate tdm	-	Checked if an existing tdm is to be downdated according to the "Document Indices" field.
Dictionary	-	Name of .mat file or workspace variable containing the dictionary to be used by tmg_query function if the "Create Query Matrix" radio button is checked.
Global Weights	-	Name of .mat file or workspace variable containing the vector of global weights to be used by tmg_query function if the "Create Query Matrix" radio button is checked.
Update Struct	-	Name of .mat file or workspace variable containing the structure to be updated or downdated by tdm_update (or tdm_downdate) function if the "Udpate tdm" or "Downdate tdm" radio button is checked.
Document Indices	-	Name of .mat file or workspace variable containing the document indices marked for deletion when the "Downdate tdm" radio button is checked.
Field Name	Default	Description

Line Delimiter	•	Checked if the “Delimiter” takes a whole line of text.
Delimiter	emptyline	The delimiter between tmg’s view of documents. Possible values are ’emptyline’, ’none_delimiter’ (treats each file as single document) or any other string.
Stoplist	-	Name of file containing stopwords, i.e. common words not used in indexing.
Min Length	3	Minimum term length.
Max Length	30	Maximum term length.
Min Local Frequency	1	Minimum local term frequency.
Max Local Frequency	inf	Maximum local term frequency.
Min Global Frequency	1	Minimum global term frequency.
Max Global Frequency	inf	Maximum global term frequency.
Local Term Weighting	TF	Local term weighting function. Possible values: ’Term Frequency’ (TF), ’Binary’, ’Logarithmic’, ’Alternate Log’, ’Augmented Normalized Term Frequency’.
Global Term Weighting	None	Global term weighting function. Possible values: ’None’, ’Entropy’, ’Inverse Document Frequency (IDF)’, ’GfIdf’, ’Normal’, ’Probabilistic Inverse’.
Database Name	-	The name of the folder (under ’data’ directory) where data are to be saved (currently supported only for the “Create New tdm” module).
Store in MySQL	-	Checked if results are to be saved into MySQL (currently supported only for the “Create New tdm” module).
use Normalization	-	Indicates normalization method. Possible values: ’None’, ’Cosine’.
use Stemming	-	Indicates if stemming is to be applied. The algorithm currrently supported is due to Porter.
Display Results	•	Display results or not to the command windows.
Continue	-	Apply the selected operation.
Reset	-	Reset window to default values.
Exit	-	Exit window.

Table 1: Description of use of `tmg_gui` components.

3.2 Dimensionality Reduction module (`dr_gui`)

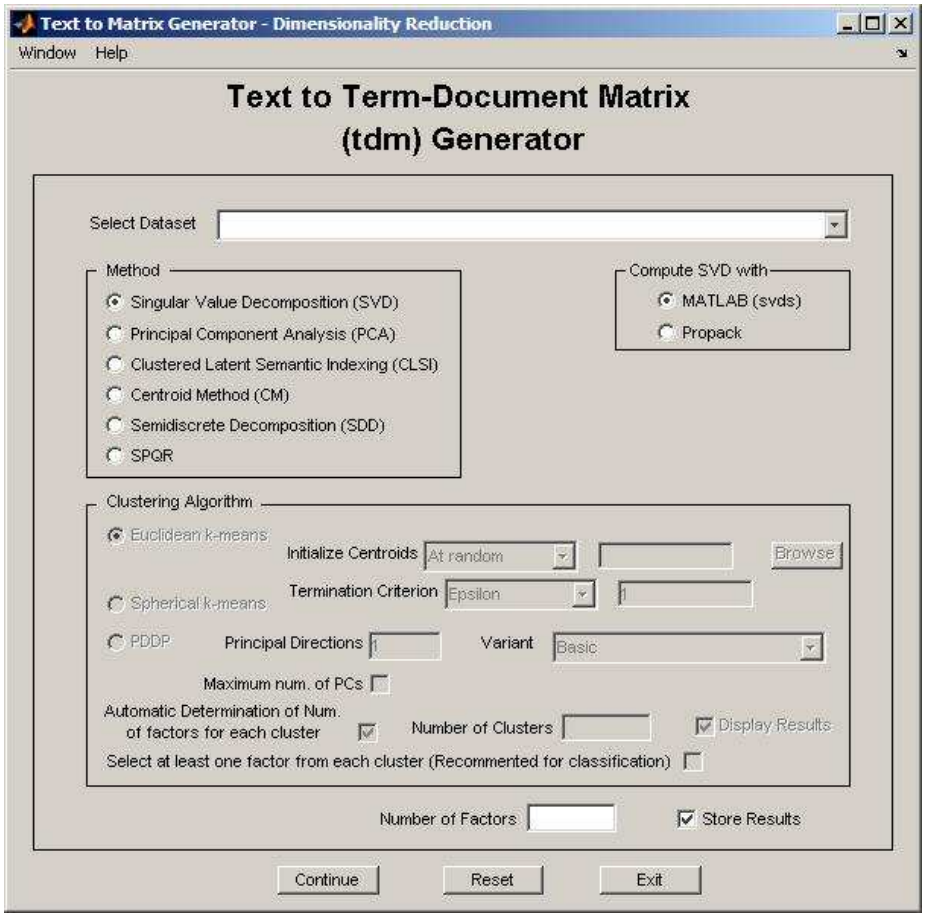

Figure 4: The `dr_gui` GUI.

This module deploys a variety of powerful techniques designed to efficiently handle high dimensional data. Dimensionality Reduction (DR) is a common technique that is widely used. The target is dual: (a) more economical representation of data, and (b) better semantic representation. TMG implements six DR techniques.

- Singular Value Decomposition (SVD).
- Principal Component Analysis (PCA).
- Clustered Latent Semantic Indexing (CLSI) [12, 13].
- Centroids Method (CM) [7].
- Semidiscrete Decomposition (SDD) [6].

- SPQR Decomposition [2].

DR data can be stored as '.mat' files and used for further processing.

The dimensionality reduction GUI module is depicted in Figure 4 while Table 2 describes in detail all the `dr_gui` fields.

Field Name	Default	Description
Select Dataset	-	Select the dataset.
Singular Value Decomposition (SVD)	•	Apply the SVD method.
Principal Component Analysis (PCA)	-	Apply the PCA method.
Clustered Latent Semantic Indexing (CLSI)	-	Apply the CLSI method.
Centroid Method (CM)	-	Apply the CM method.
Semidiscrete Decomposition (SDD)	-	Apply the SDD method.
SPQR	-	Apply the SPQR method.
MATLAB (svds)	•	Check to use MATLAB function svds for the computation of the SVD or PCA.
Propack	-	Check to use PROPACK package for the computation of the SVD or PCA.
Euclidean k-means	•	Check to use the euclidean k-means clustering algorithm in the course of CLSI or CM.
Spherical k-means	-	Check to use the spherical k-means clustering algorithm in the course of CLSI or CM.
PDDP	-	Check to use the PDDP clustering algorithm in the course of CLSI or CM.
Initialize Centroids	At random	Defines the method used for the initialization of the centroid vector in the course of k-means. Possibilities are: initialize at random and supplly a variable of '.mat' file with the centroids matrix.
Termination Criterion	Epsilon (1)	Defines the termination criterion used in the course of k-means. Possibilities are: use an epsilon value (default 1) and stop iteration when the objective function improvement does not exceed epsilon or perform a specific number of iterations (default 10).
Principal Directions	1	Number of principal directions used in PDDP.

Maximum num. of PCs	-	Check if the PDDP(max-l) variant is to be applied.
Variant	Basic	A set of PDDP variants. Possibe values: 'Basic', 'Split with k-means', 'Optimat Split', 'Optimal Split with k-means', 'Optimal Split on Projection'.
Automatic Determination of Num. of factors for each cluster	•	Check to apply a heuristic for the determination of the number of factors computed from each cluster in the course of the CLSI algorithm.
Number of Clusters	-	Number of clusters computed in the course of the CLSI algorithm.
Display Results	•	Display results or not to the command windows.
Select at least one factor from each cluster	-	Use this option in case low-rank data are to be used in the course of classification.
Number of factors	-	Rank of approximation.
Store Results	•	Check to store results.
Continue	-	Apply the selected operation.
Reset	-	Reset window to default values.
Exit	-	Exit window.

Table 2: Description of use of `dr_gui` components.

3.3 Retrieval module (`retrieval_gui`)

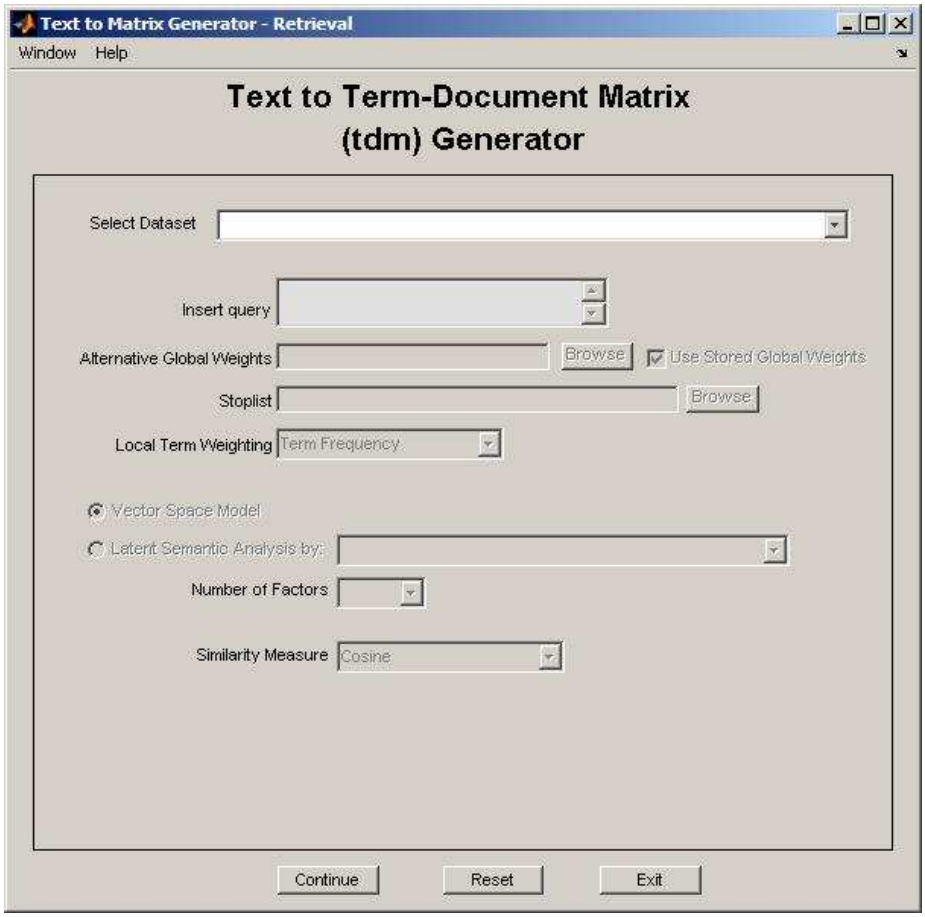

Figure 5: The `retrieval_gui` GUI.

TMG offers two alternatives for Text Mining.

- Vector Space Model (VSM) [9].
- Latent Semantic Analysis (LSA) [1, 4],

using a combination of any DR technique and Latent Semantic Indexing (LSI). Using the corresponding GUI, the user can apply a question to an existing dataset using any of the aforementioned techniques and get HTML response.

The retrieval GUI module is depicted in Figure 5 while Table 3 describes in detail all the `retrieval_gui` fields.

Field Name	Default	Description
Select Dataset	-	Select the dataset.
Insert Query	•	The query to be executed.
Alternative Global Weights	-	Global weights vector used for the construction of the query vector.
Use Stored Global Weights	•	Use the global weights vector found on the container directory of the dataset.
Stoplist	-	Use a stoplist.
Local Term Weighting	TF	The local term weighting to be used.
Vector Space Model	•	Apply the Vector space Model retrieval method.
Latent Semantic Analysis	-	The method used in the course of the Latent Semantic Analysis technique. Possible values: 'Singular Value Decomposition', 'Principal Component Analysis', 'Clustered Latent Semantic Analysis', 'Centroid Mathod', 'Semidiscrete Decomposition', 'SPQR'.
Number of Factors	-	Select the number of factors used during the retrieval process.
Similarity Measure	Cosine	Similarity measure used during the retrieval process.
Continue	-	Apply the selected operation.
Reset	-	Reset window to default values.
Exit	-	Exit window.

Table 3: Description of use of `retrieval_gui` components.

3.4 Clustering module (`clustering_gui`)

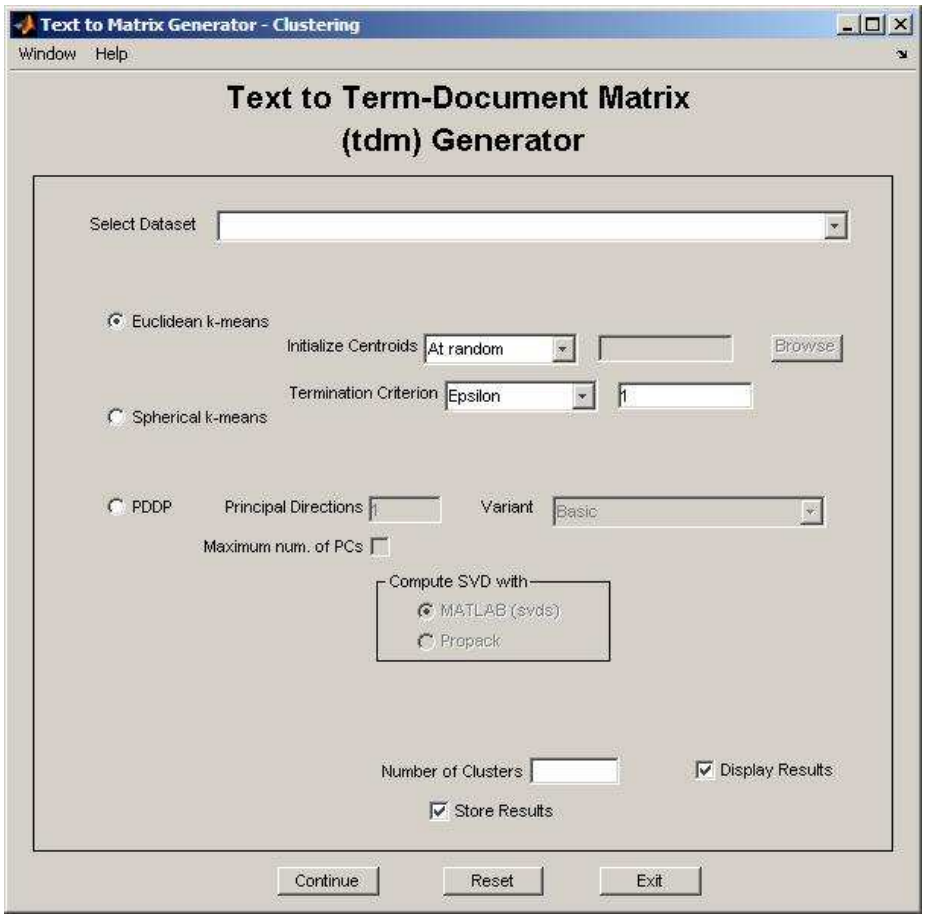

Figure 6: The `clustering_gui` GUI.

TMG implements three clustering algorithms.

- k-means.
- Spherical k-means [5].
- Principal Direction Divisive Partitioning (PDDP) [3, 11].

Regarding PDDP, TMG implements the basic algorithm as well as the PDDP(l) [11] along with some recent hybrid variants of PDDP and kmeans [15].

The clustering GUI module is depicted in Figure 6 while Table 4 describes in detail all the `clustering_gui` fields.

Field Name	Default	Description
Select Dataset	-	Select the dataset.
Euclidean k-means	•	Check to use the euclidean k-means clustering algorithm.
Spherical k-means	-	Check to use the spherical k-means clustering algorithm.
PDDP	-	Check to use the PDDP clustering algorithm.
Initialize Centroids	At random	Defines the method used for the initialization of the centroid vector in the course of k-means. Possibilities are: initialize at random and supplly a variable of '.mat' file with the centroids matrix.
Termination Criterion	Epsilon (1)	Defines the termination criterion used in the course of k-means. Possibilities are: use an epsilon value (default 1) and stop iteration when the objective function improvement does not exceed epsilon or perform a specific number of iterations (default 10).
Principal Directions	1	Number of principal directions used in PDDP.
Maximum num. of PCs	-	Check if the PDDP(max-l) variant is to be applied.
Variant	Basic	A set of PDDP variants. Possibe values: 'Basic', 'Split with k-means', 'Optimat Split', 'Optimal Split with k-means', 'Optimal Split on Projection'.
MATLAB (svds)	•	Check to use MATLAB function svds for the computation of the SVD in the course of PDDP.
Propack	-	Check to use PROPACK package for the computation of the SVD in the course of PDDP.
Number of Clusters	-	Number of clusters computed.
Display Results	•	Display results or not to the command windows.
Store Results	•	Check to store results.
Continue	-	Apply the selected operation.
Reset	-	Reset window to default values.
Exit	-	Exit window.

Table 4: Description of use of `clustering_gui` components.

3.5 Classification module (`classification_gui`)

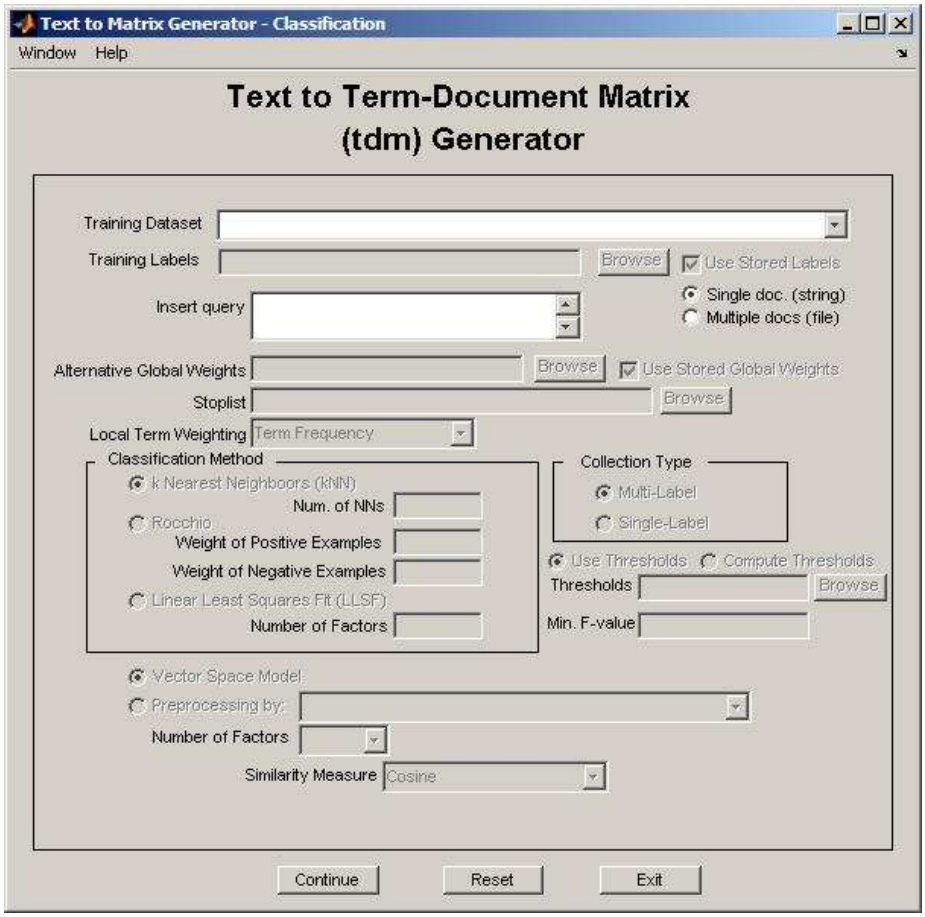

Figure 7: The `classification_gui` GUI.

TMG implements three classification algorithms.

- k Nearest Neighboors (kNN).
- Rocchio.
- Linear Least Squares Fit (LLSF) [10].

All these algorithms can be combined with CLSI, CM and SVD DR techniques.

The classification GUI module is depicted in Figure 7 while Table 5 describes in detail all the `classification_gui` fields.

Field Name	Default	Description
Training Dataset	-	The training dataset.
Training Labels	-	The labels of the training dataset.
Use Stored Labels	•	Check to use the stored vector of labels of training documents in the container folder.
Insert query	-	The test document(s).
Single doc. (string)	•	Check if a single test document is to be inserted.
Multiple docs (file)	-	Check if multiple test document are to be inserted.
Filename	-	In 'Multiple docs (file)' is checked, insert the filename containing the test documents.
Delimiter	-	In 'Multiple docs (file)' is checked, insert the delimiter o be used for the test documents.
Line Delimiter	•	In 'Multiple docs (file)' is checked, check if delimiter of test documents' file takes a whole l of text.
Alternative Global Weights	-	Global weights vector used for the construction of the test documents' vectors.
Use Stored Global Weights	•	Use the global weights vector found on the container directory of the training dataset.
Stoplist	-	Use a stoplist.
Local Term Weighting	TF	The local term weighting to be used.
k Nearest Neighboors (kNN)	•	Check if the kNN classifier is to be applied.
Num. of NNs	-	Number of Nearest Neighboors in kNN classifier.
Rocchio	-	Check if Rocchio classifier is to be applied.
Weight of Positive Examples	-	The weight of the positive examples in the formation of the centroids vectors in Rocchio.
Weight of Negative Examples	-	The weight of the negative examples in the formation of the centroids vectors in Rocchio.
Linear Least Squares Fit (LLSF)	-	Check if LLSF classifier is to be applied.
Number of Factors	-	Number of factors used in the course of LLSF.
Multi-Label	•	Check if classifier is to be applied for a multi-label collection.
Single-Label	-	Check is classifier is to be applied for a single-label collection.
Use Thresholds	•	If 'Multi-Label' radio button is checked, use a stored vector of thresholds.
Compute Thresholds	-	If 'Multi-Label' radio button is checked, compute thresholds.
Thresholds	-	If 'Multi-Label' and 'Use Thresholds' radio buttons are checked, supply a stored vector of thresholds.

Min. F-value	-	If 'Multi-Label' and 'Compute Thresholds' radio buttons are checked, supply minimum F1 value used in the thresholding algorithm.
Vector Space Model	•	Use the basic Vector Space Model.
Preprocessing by	-	Use preprocessed training data with: 'Singular Value Decomposition', 'Principal Component Analysis', 'Clusteredd Latent Semantic Analysis', 'Centroid Mathod', 'Semidiscrete Decomposition', 'SPQR'.
Number of Factors	-	Number of factors for preprocessed training data.
Similarity Measure	Cosine	The similarity measure to be used.
Continue	-	Apply the selected operation.
Reset	-	Reset window to default values.
Exit	-	Exit window.

Table 5: Description of use of `classification_gui` components.

Acknowledgments

TMG was conceived after a motivating discussion with Andrew Knyazev regarding a collection of MATLAB tools we had put together to aid in our clustering experiments. We thank our collegues Ioannis Antonellis, Anastasios Zouzias, Efi Kokiopoulou and Constantine Bekas for many helpful suggestions, Jacob Kogan and Charles Nicholas for inviting us to contribute to [14], Elias Houstis for his help in the initial phases of this research and Michael Berry, Tamara Kolda and Rasmus Munk Larsen for letting us use and distribute SPQR, SDDPACK and PROPACK software respectively. Special thanks are due to many of the users for their constructive comments regarding TMG. This research was supported in part by a University of Patras "Karatheodori" grant. The first author was also supported by a Bodossaki Foundation graduate fellowship.

References

[1] M. Berry, Z. Drmac, and E. Jessup, *Matrices, vector spaces, and information retrieval*, SIAM Review **41** (1998), 335–362.

[2] M. W. Berry, S. A. Pulatova, and G. W. Stewart, *Computing sparse reduced-rank approximations to sparse matrices*, ACM TOMS **31** (2005), no. 2.

[3] D. Boley, *Principal direction divisive partitioning*, Data Mining and Knowledge Discovery **2** (1998), no. 4, 325–344.

[4] S. Deerwester, S. Dumais, G. Furnas, T. Landauer, and Harshman R., *Indexing by Latent Semantic Analysis*, Journal of the American Society for Information Science **41** (1990), no. 6, 391–407.

[5] I. S. Dhillon and D. S. Modha, *Concept decompositions for large sparse text data using clustering*, Machine Learning **42** (2001), no. 1, 143–175.

[6] T. Kolda and D. O'Leary, *Algorithm 805: computation and uses of the semidiscrete matrix decomposition*, ACM TOMS **26** (2000), no. 3.

[7] H. Park, M. Jeon, and J. Rosen, *Lower dimensional representation of text data based on centroids and least squares*, BIT **43** (2003).

[8] M.F. Porter, *An algorithm for suffix stripping*, Program (1980), no. 3, 130–137.

[9] G. Salton, C. Yang, and A. Wong, *A Vector-Space Model for Automatic Indexing*, Communications of the ACM **18** (1975), no. 11, 613–620.

[10] Y. Yang and C. Chute, *A linear least squares fit mapping method for information retrieval from natural language texts*, In 14th Conf. Comp. Linguistics, 1992.

[11] D. Zeimpekis and E. Gallopoulos, *PDDP(l): Towards a Flexing Principal Direction Divisive Partitioning Clustering Algorithms*, Proc. IEEE ICDM '03 Workshop on Clustering Large Data Sets (Melbourne, Florida) (D. Boley, I. Dhillon, J. Ghosh, and J. Kogan, eds.), 2003, pp. 26–35.

[12] D. Zeimpekis and E. Gallopoulos, *CLSI: A flexible approximation scheme from clustered term-document matrices*, In Proc. SIAM 2005 Data Mining Conf. (Newport Beach, California) (H. Kargupta, J. Srivastava, C. Kamath, and A. Goodman, eds.), April 2005, pp. 631–635.

[13] D. Zeimpekis and E. Gallopoulos, *Linear and non-linear dimensional reduction via class representatives for text classification*, In Proc. of the 2006 IEEE International Conference on Data Mining (Hong Kong), December 2006, pp. 1172–1177.

[14] D. Zeimpekis and E. Gallopoulos, *TMG: A MATLAB toolbox for generating term-document matrices from text collections*, Grouping Multidimensional Data: Recent Advances in Clustering (J. Kogan, C. Nicholas, and M. Teboulle, eds.), Springer, Berlin, 2006, pp. 187–210.

[15] D. Zeimpekis and E. Gallopoulos, *k-means steering of spectral divisive clustering algorithms*, In Proc. of Text Mining Workshop (Minneapolis), 2007.

A Appendix: Demonstration of Use

A.1 Indexing module (`tmg_gui`)

Assume we want to run `tmg.m` for the following input:

- filename: sample_documents/sample1
- delimiter: emptyline
- line_delimiter: yes
- stoplist: common_words
- minimum length: 3
- maximum length: 30
- minimum local frequency: 1
- maximum local frequency: inf
- minimum global frequency: 1
- maximum global frequency: inf
- local term weighting: logarithmic
- global term weighting: IDF
- normalization: cosine
- stemming: -

and you want to store results in sample1 directory and in MySQL.

1. Initially select the operation you want to perform, by pressing the corresponding radio button to the upper frame.

2. The selection of a radio button activates the required fields in the GUI, while deactivating the rest fields and changing their background color.

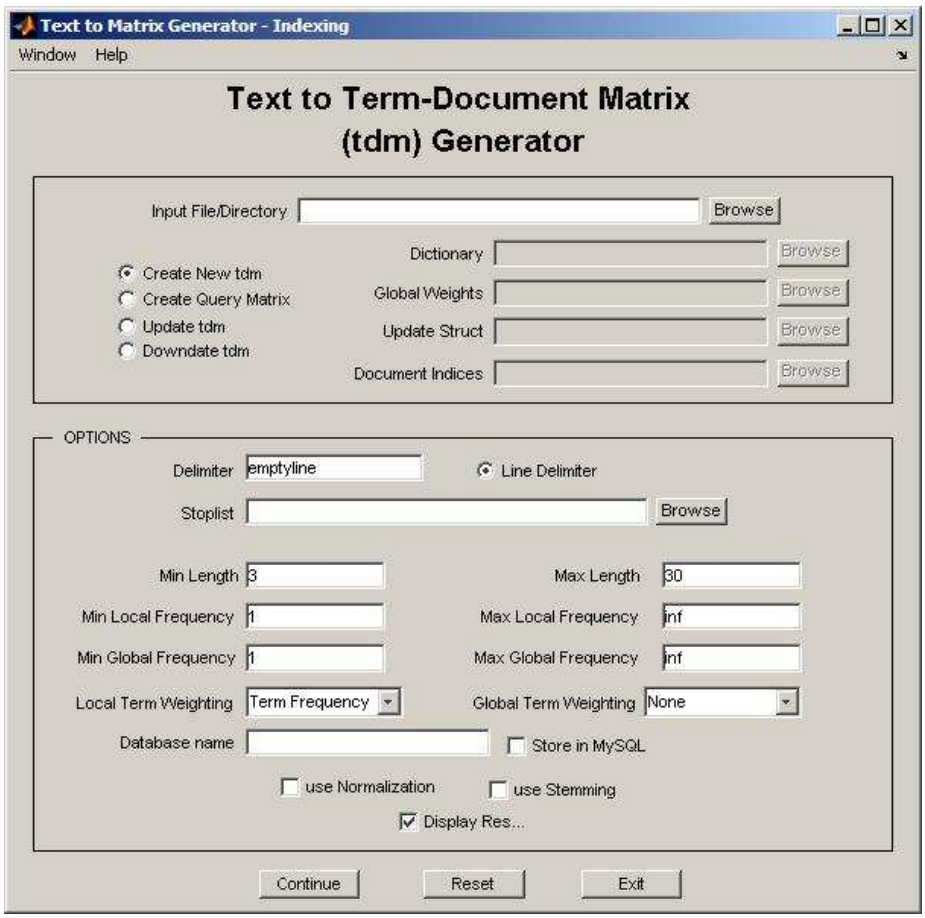

Figure 8: Starting window of `tmg_gui`.

3. Fill in the required fields, by pressing the check buttons, editing the edit boxes or selecting the appropriate files/variables by pressing a “Browse” button.

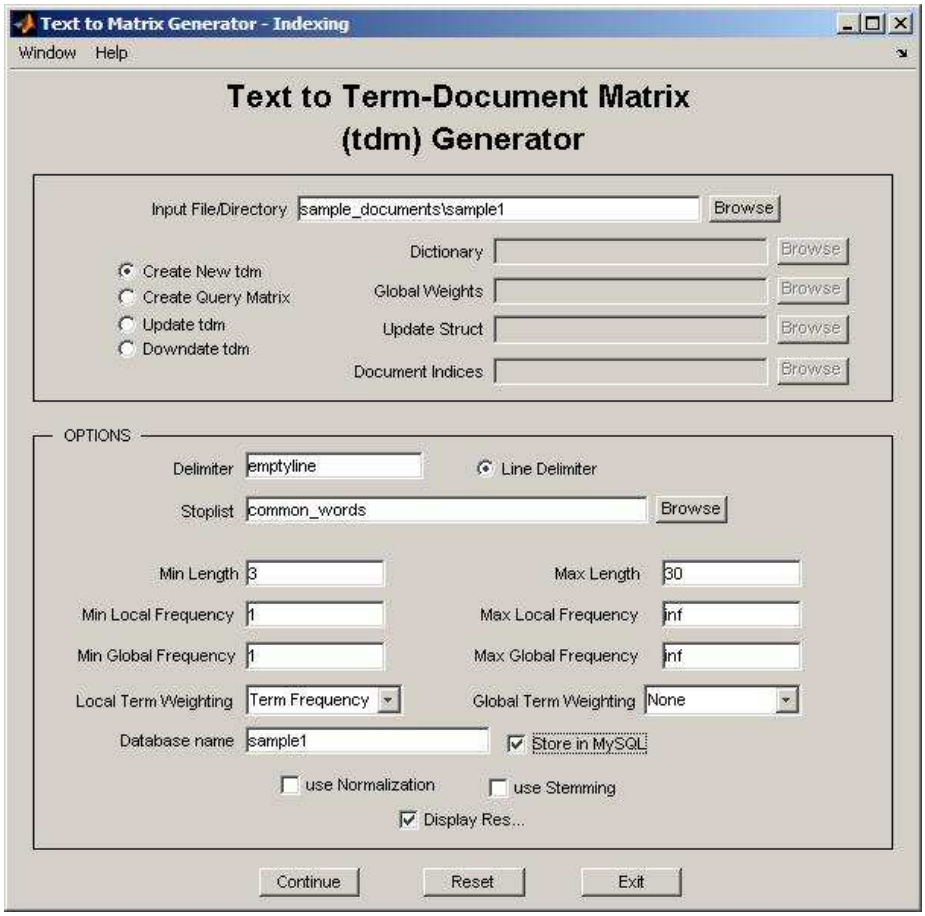

Figure 9: Next view of `tmg_gui` according to the user selection.

4. Pressing a “Browse” button the user has the ability to chose a file or a variable.

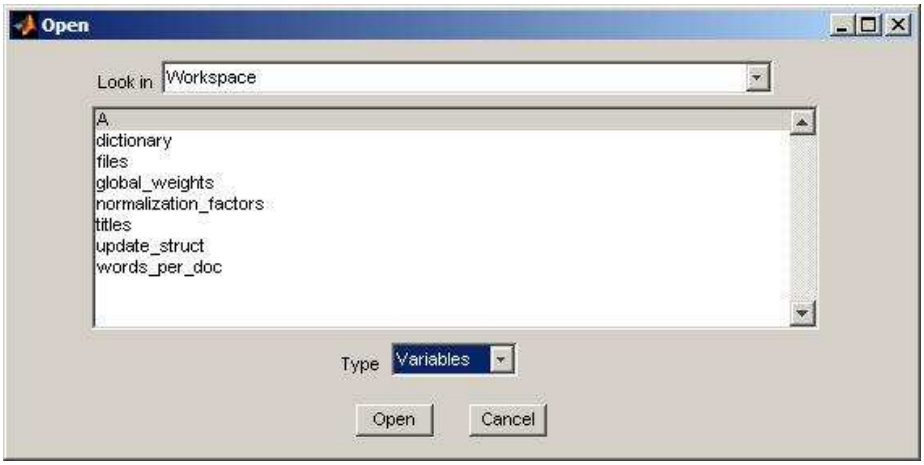

Figure 10: The `open_file` window.

5. Press the ”Continue” button in order to perform the selected operation.

6. Results have been saved to the workspace. Furthermore, directory ‘sample1’ has been created under ‘TMG_HOME/data’ with each output variable stored in a single ‘.mat’ file.

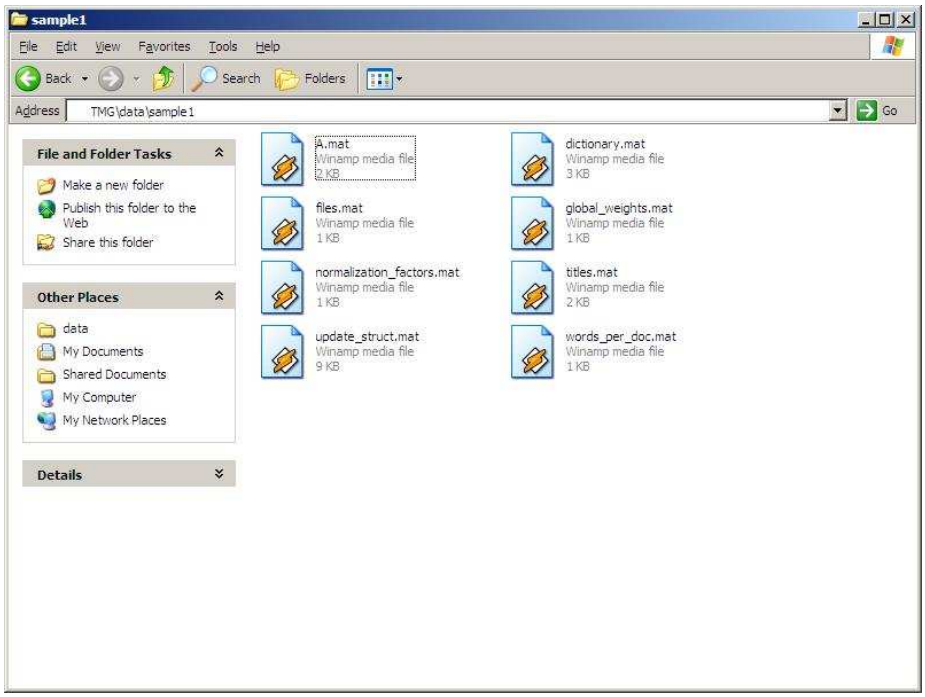

Figure 11: The output ‘.mat’ files of `tmg_gui`.

7. Results have also been saved in MySQL (used for further processing, e.g. `retrieval_gui`).

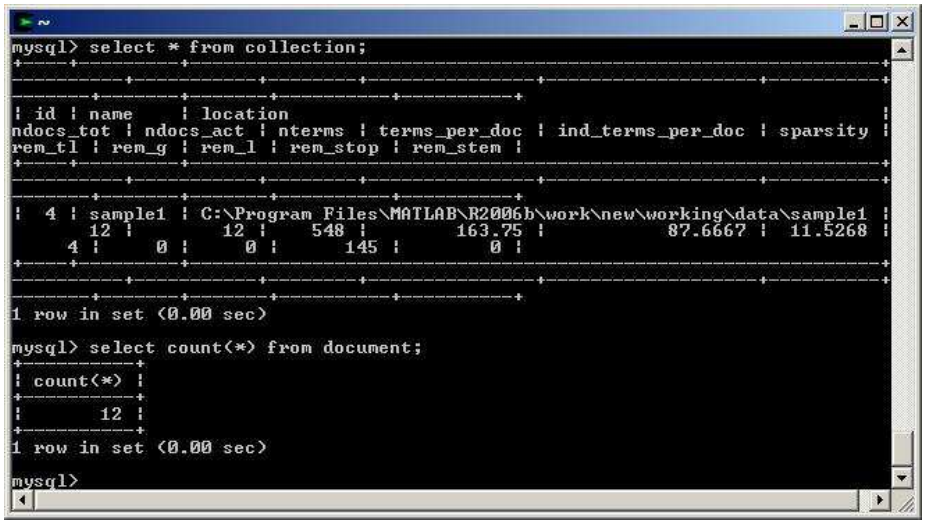

Figure 12: The MySQL view uppon `tmg` execution.

8. Press the “Reset” button in order to change the input.

9. For further documentation type '`help tmg_gui`' at the MATLAB command window, or select the "Documentation" tab from the "Help" menu.

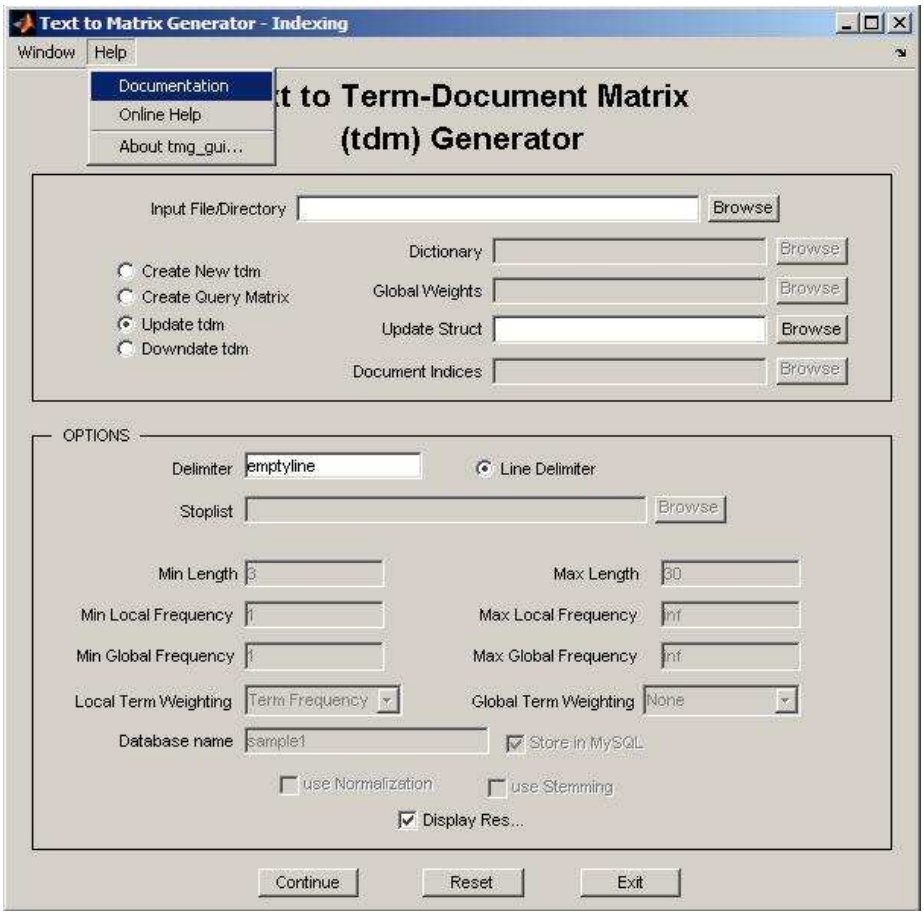

Figure 13: The GUIs' general help tab.

10. In order to update a tdm, give the "input file/directory" and the update_struct corresponding to the initial collection. In case you just want to alter some options, give a blank "input file/direcory" and change the corresponding fields of update_struct.

11. In order to downdate a tdm, give the update_struct corresponding to the initial collection and the document indices vector you want to remove.

12. In order to construct a term-query matrix, give the dictionary char array of the initial collection and the corresponding vector of global weights (optional).

A.2 Dimensionality Reduction module (`dr_gui`)

Suppose we have processed a collection with `tmg_gui`, construct a tdm with 1,033 documents and 12,184 terms (corresponding to the well-known MEDLINE collection) and store the results in 'TMG_HOME/data/medline'. Assume then, we want to construct a low-rank approximation of the tdm ,using the Clustered Latent Semantic Indexing (CLSI) technique for the following input:

- compute SVD with: Propack
- clustering algorithm: PDDP
- principal directions: 1
- maximum number of PCs: -
- variant: basic
- automatic determination of num. of factors from each cluster: yes
- number of clusters: 10
- number of factors: 100

and you want to store results in the medline directory.

1. Initially select the operation you want to perform, by pressing the corresponding radio button to the upper left frame.

2. The selection of a radio button activates the required fields in the GUI, while deactivating the rest fields and changing their background color.

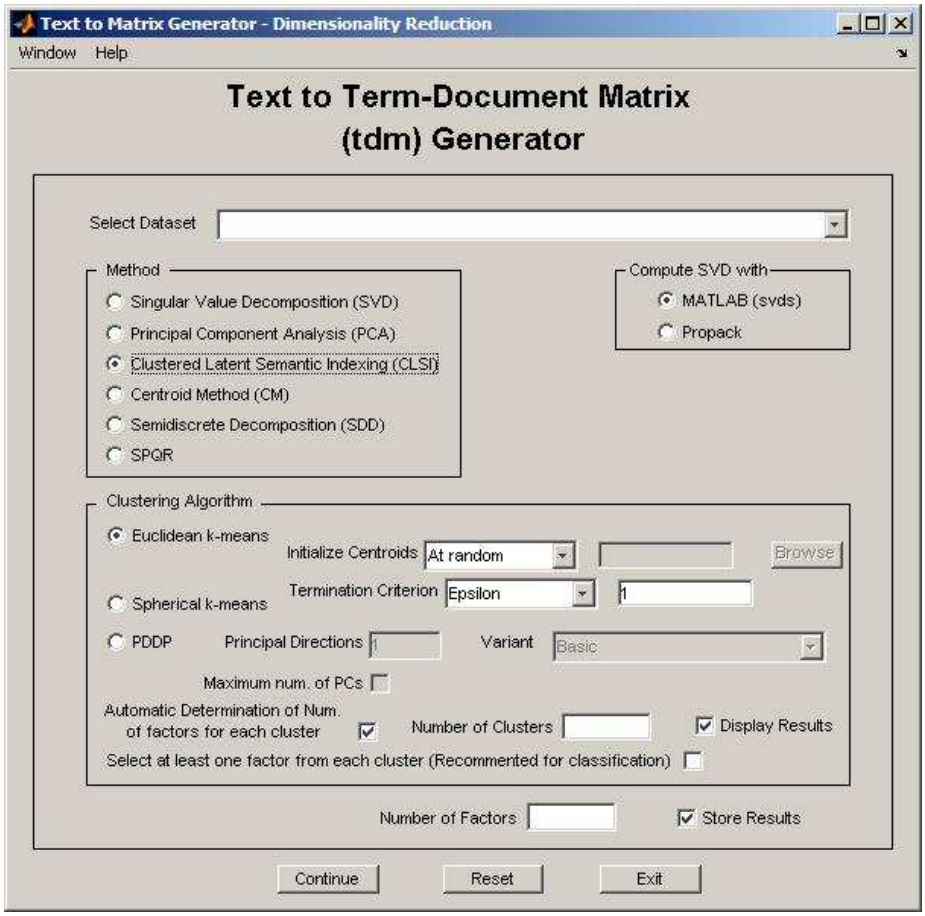

Figure 14: Starting window of `dr_gui`.

3. Fill in the required fields, by pressing the check buttons, editing the edit boxes or selecting the appropriate files/variables by pressing a “Browse” button.

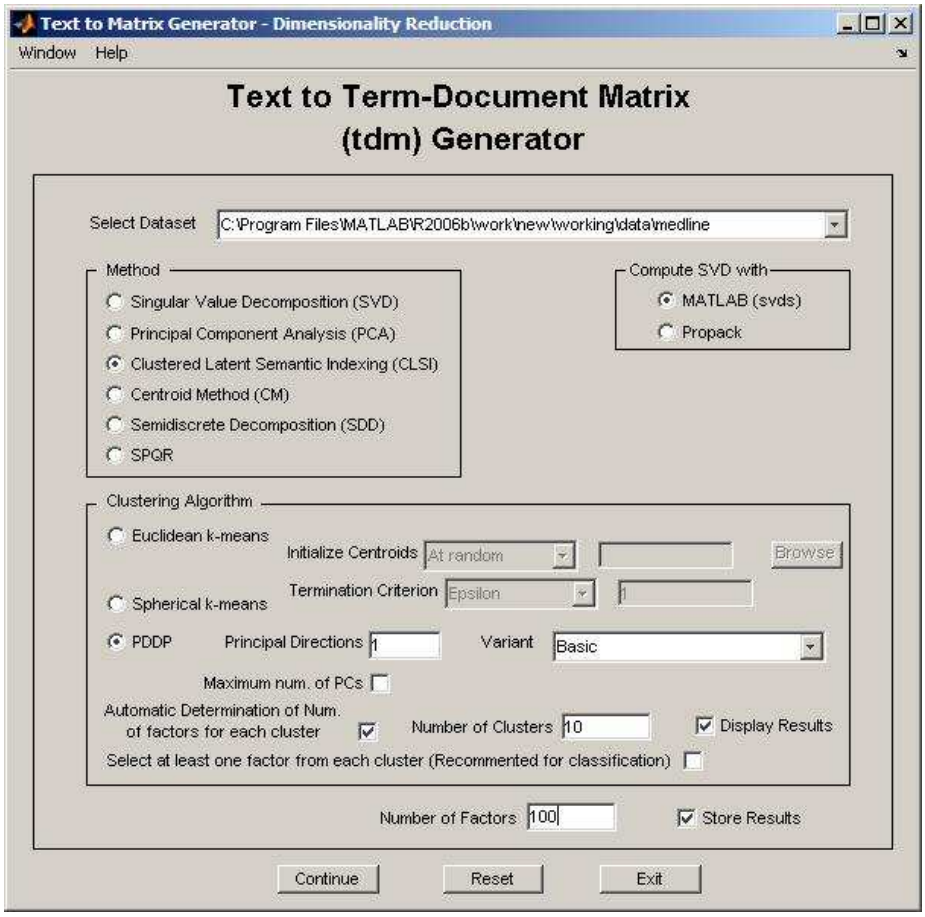

Figure 15: Next view of `dr_gui` according to the user selection.

4. Press the “Continue” button in order to perform selected operation.

5. Results have been saved to the workspace. Furthermore, directory ‘clsi/k_100’ has been created under ‘TMG_HOME/data/medline’ with each output variable stored in a single ‘.mat’ file.

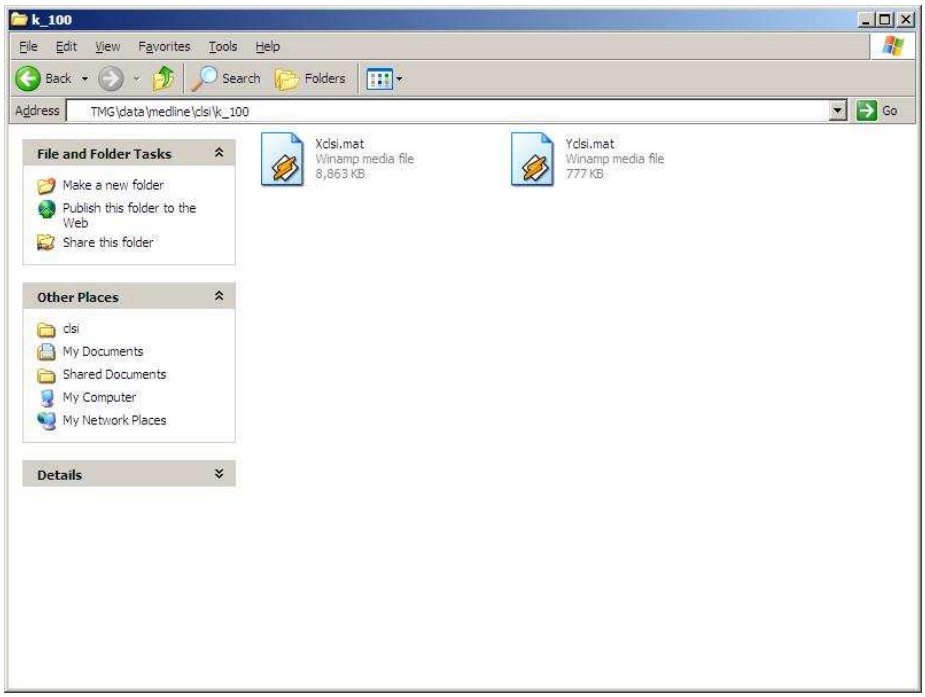

Figure 16: The output ‘.mat’ files of `dr_gui`.

6. Press the “Reset” button in order to change the input.

A.3 Retrieval module (`retrieval_gui`)

Suppose we have processed a collection `with tmg_gui`, construct a tdm with 1,033 documents and 12,184 terms (corresponding to the well-known MEDLINE collection) and store the results in 'TMG_HOME/data/medline'. Assume then, we want to retrieve the relevant documents to a specific query for the following input:

- insert query: 'the crystalline lens in vertebrates, including humans'
- use stored global weights: yes
- stoplist: common_words
- local term weighting: Term Frequency
- latent semantic analysis by: Clustered Latent Semantic Indexing
- number of factors: 100
- similarity measure: Cosine

1. Initially select the retrieval method you want to apply, by pressing the corresponding radio button.

2. The selection of a radio button activates the required fields in the GUI, while deactivating the rest fields and changing their background color.

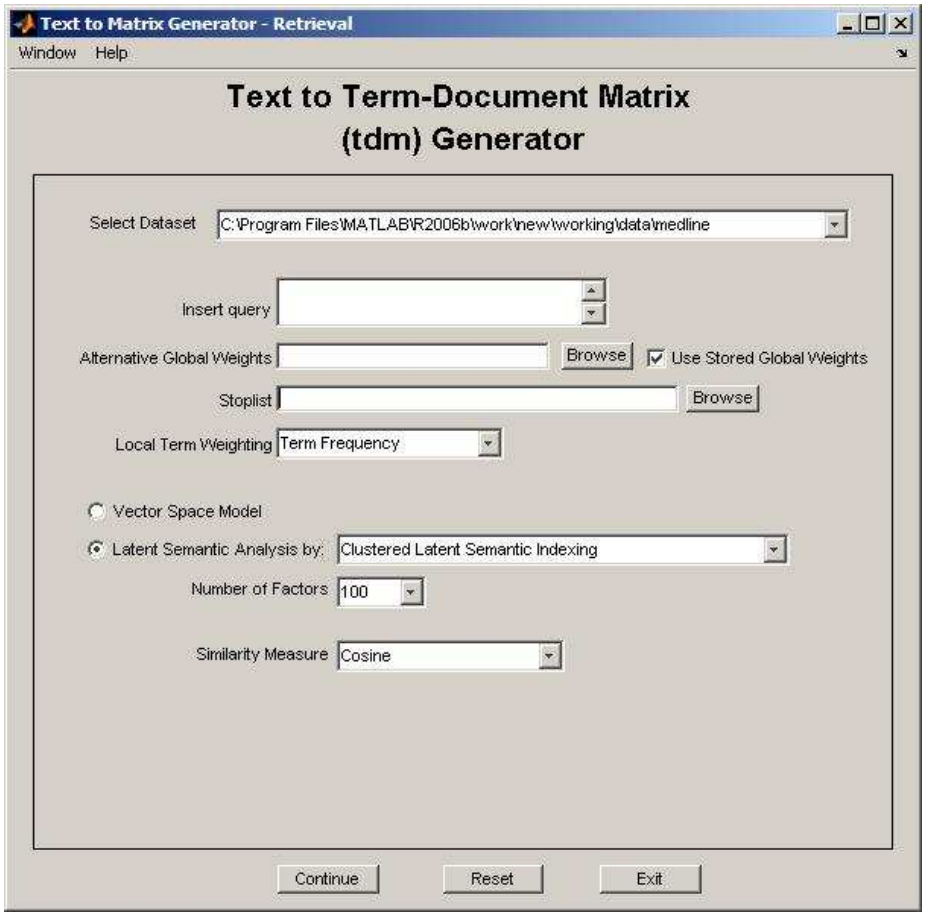

Figure 17: Starting window of `retrieval_gui`.

3. Fill in the required fields, by pressing the check buttons, editing the edit boxes or selecting the appropriate files/variables by pressing a “Browse” button.

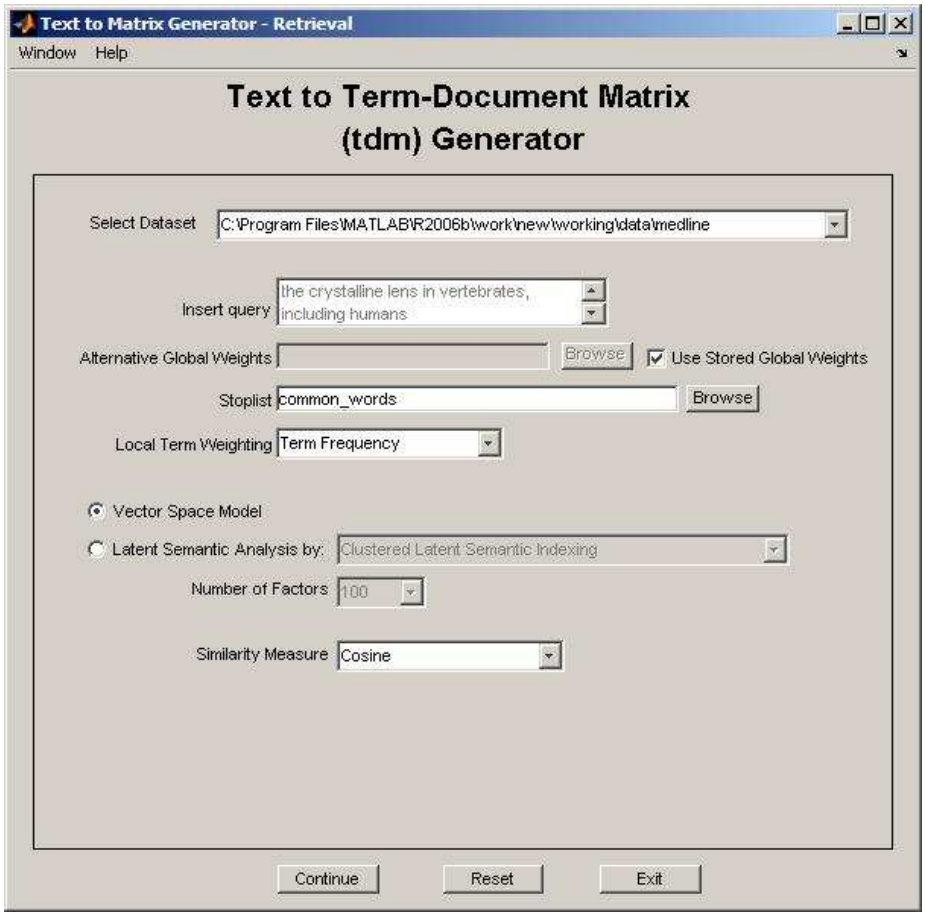

Figure 18: Next view of `retrieval_gui` according to the user selection.

4. Press the “Continue” button in order to perform selected operation.
5. Results have been saved to the workspace.
6. Furthermore, in case data have been stored to MySQL, the user gets an html response.

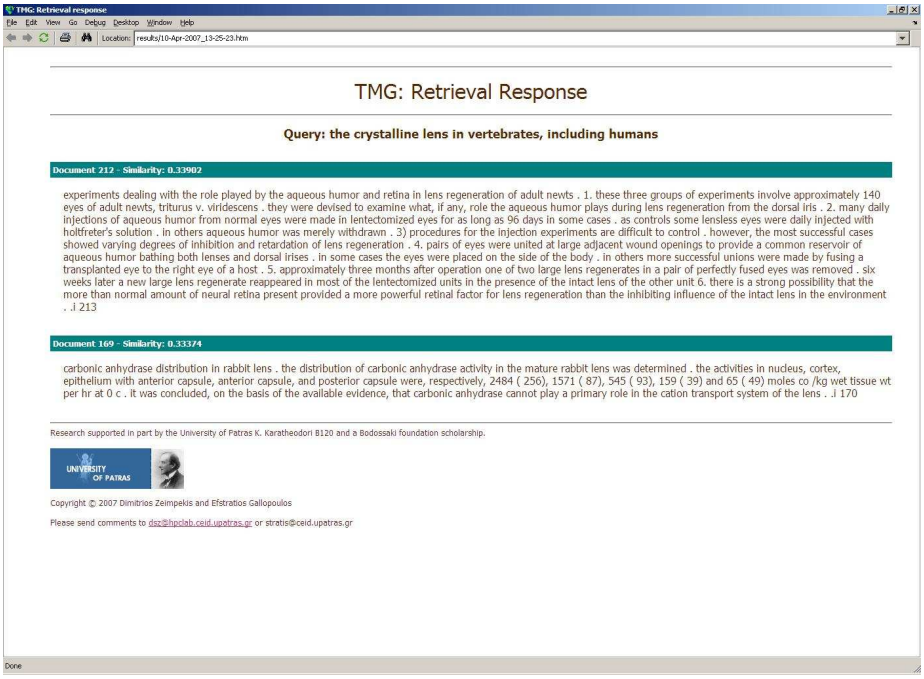

Figure 19: The output of `retrieval_gui`.

7. Press the “Reset” button in order to change the input.

A.4 Clustering module (`clustering_gui`)

Suppose we have processed a collection with `tmg_gui`, construct a tdm with 1,033 documents and 12,184 terms (corresponding to the well-known MEDLINE collection) and store the results in 'TMG_HOME/data/medline'. Assume then, we want to cluster the tdm ,using the k-means clustering algorithm with the following input:

- initialize centroids: At random
- termination criterion: Num. iterations, value 10
- number of clusters: 10

and you want to store results in the medline directory.

1. Initially select the clustering algorithm you want to apply, by pressing the corresponding radio button.

2. The selection of a radio button activates the required fields in the GUI, while deactivating the rest fields and changing their background color.

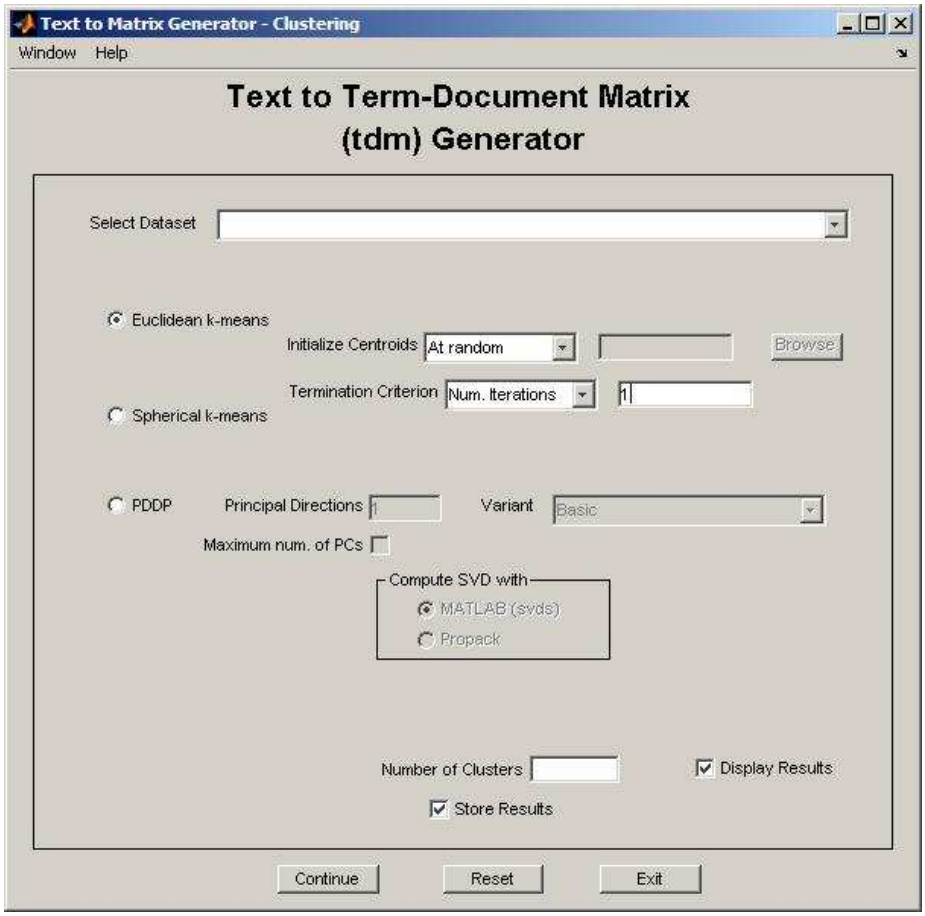

Figure 20: Starting window of `clustering_gui`.

3. Fill in the required fields, by pressing the check buttons, editing the edit boxes or selecting the appropriate files/variables by pressing a “Browse” button.

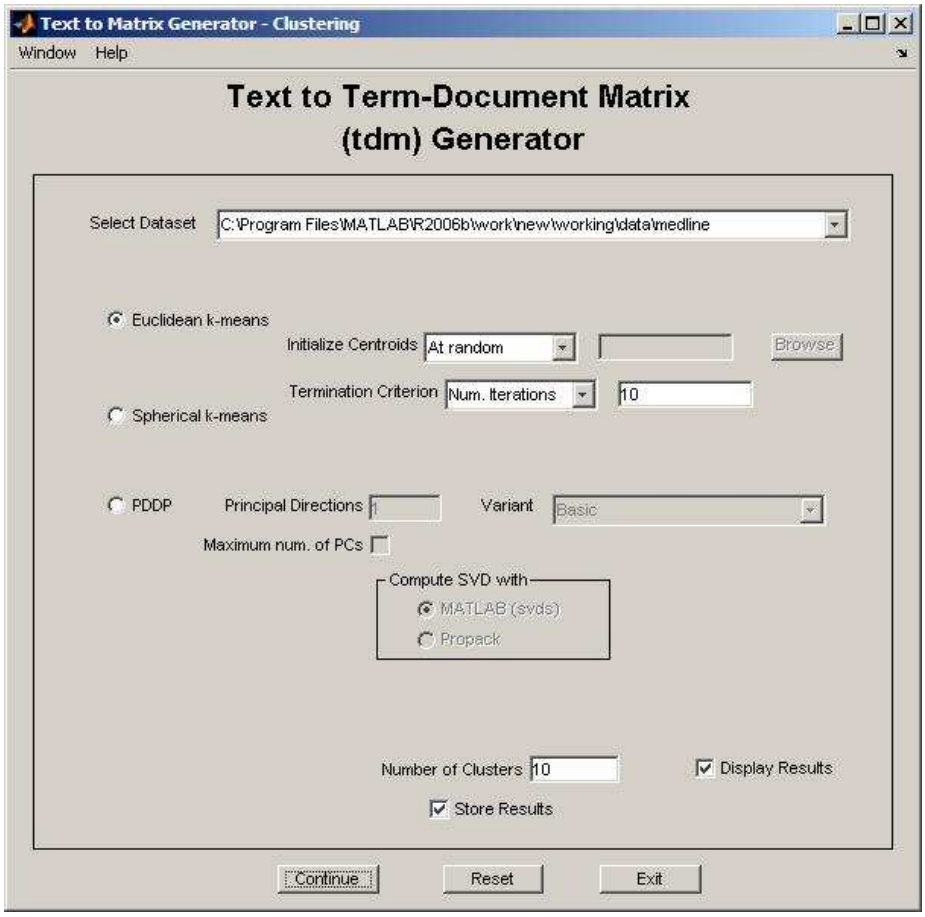

Figure 21: Next view of `clustering_gui` according to the user selection.

4. Press the “Continue” button in order to perform selected operation.

5. Results have been saved to the workspace. Furthermore, directory ‘kmeans/k_10’ has been created under ‘TMG_HOME/data/medline’ with each output variable stored in a single ‘.mat’ file.

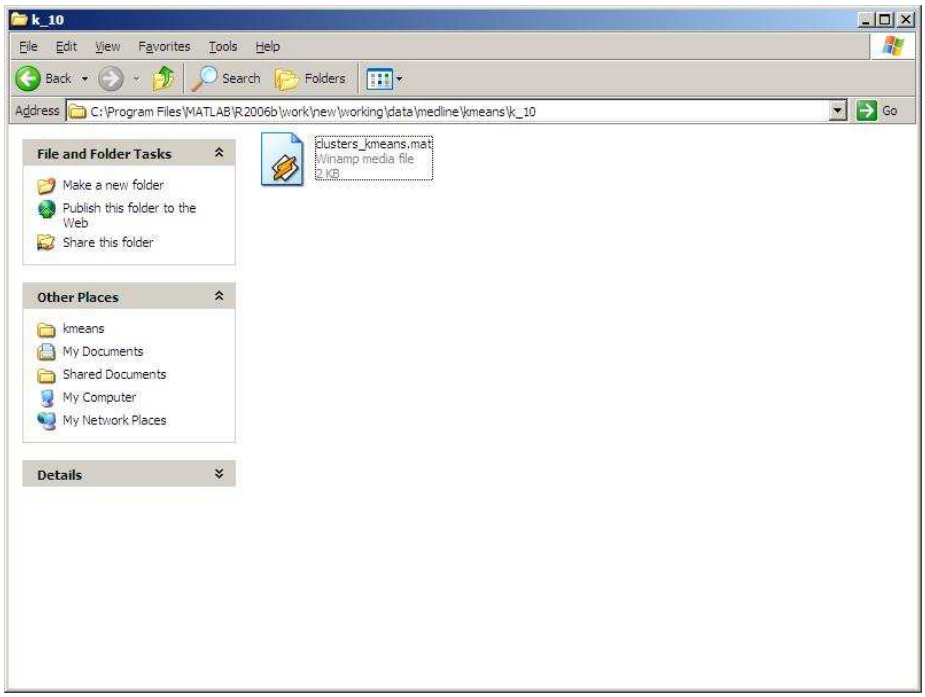

Figure 22: The output ‘.mat’ files of `clustering_gui`.

6. The user gets an html response that summarizes the clustering result.

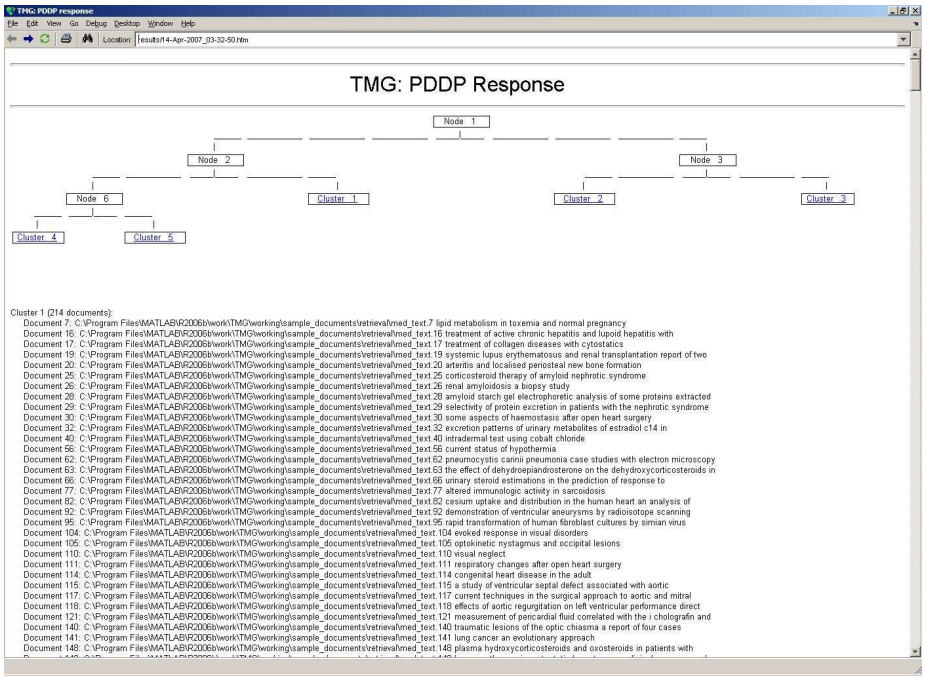

Figure 23: The output of `clustering_gui` for PDDP.

7. Press the “Reset” button in order to change the input.

A.5 Classification module (`classification_gui`)

Suppose we have processed a collection with `tmg_gui`, construct a tdm with 6,495 documents and 21,764 terms (a single label dataset corresponding to the well-known modapte split of the Reuters-21578 collection) and store the results in 'TMG_HOME/data/reuters'. Assume then, we want to classify the test part of the modapte split,using the k-Nearest Neighboors classifier for the following input:

- Multiple docs (file): yes
- filename: sample_document/reuters.test
- delimiter: </reuters>
- line delimiter: yes
- use stored global weights: yes
- stoplist: common_words
- local term weighting: Term Frequency
- classification method: k Nearest Neighboors (kNN)
- num. of NNs: 10
- collection type: Single-Label
- preprocessed by: Clustered Latent Semantic Indexing
- number of factors: 100
- similarity measure: Cosine

1. Initially select the classification algorithm you want to apply, by pressing the corresponding radio button to left frame.

2. The selection of a radio button activates the required fields in the GUI, while deactivating the rest fields and changing their background color.

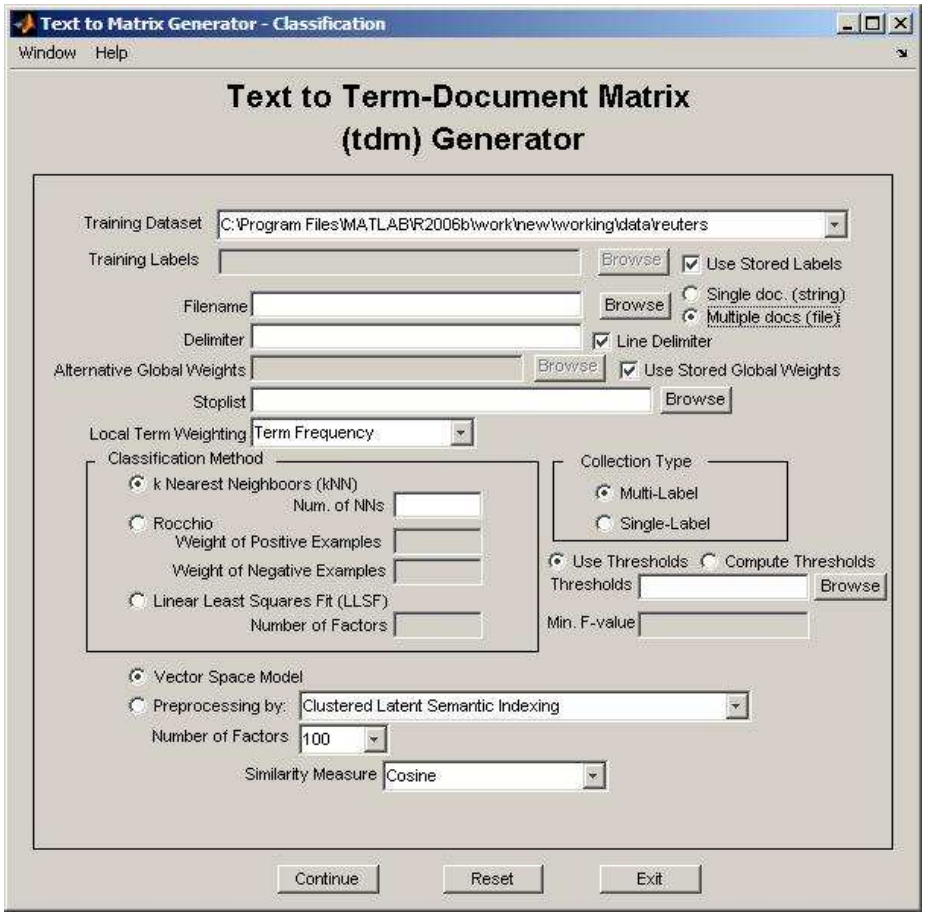

Figure 24: Starting window of `classification_gui`.

3. Fill in the required fields, by pressing the check buttons, editing the edit boxes or selecting the appropriate files/variables by pressing a “Browse” button.

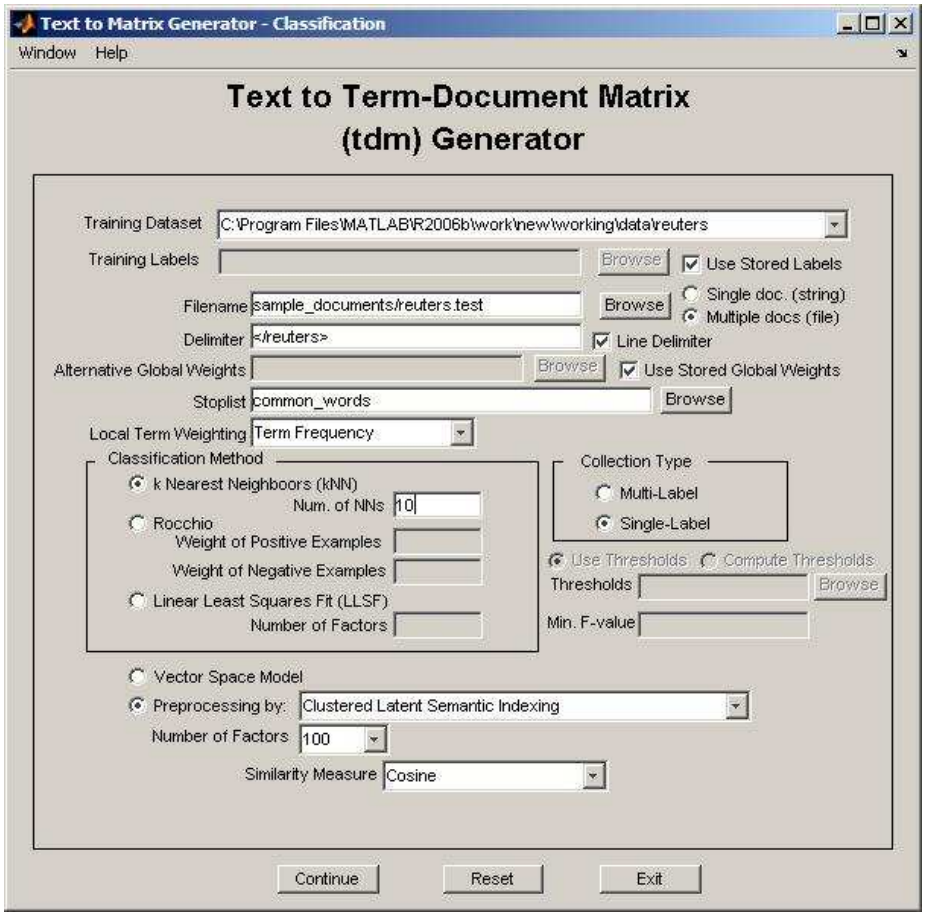

Figure 25: Next view of `classification_gui` according to the user selection.

4. Press the “Continue” button in order to perform selected operation.
5. Results have been saved to the workspace.
6. Press the “Reset” button in order to change the input.

B Appendix: Function Reference

about_tmg_gui
ABOUT_TMG_GUI ABOUT_TMG_GUI displays information for TMG.

block_diagonalize
BLOCK_DIAGONALIZE - reorders a matrix heuristically using a clustering result [A, N_ROWS, N_COLS, ROW_INDS, COL_INDS]=BLOCK_DIAGONALIZE(A, CLUSTERS) reorders matrix A using the clustering result represented by the structure CLUSTERS. N_ROWS and N_COLS store the last row and column index for each row and column block resprectively, while ROW_INDS and COL_INDS contain the permuted row and column indices.

classification_gui
CLASSIFICATION_GUI CLASSIFICATION_GUI is a graphical user interface for all classification functions of the Text to Matrix Generator (TMG) Toolbox.

`clsi`

CLSI - computes a rank-L approximation of the input matrix
using the Clustered Latent Semantic Indexing Method [1]

[X, Y]=CLSI(A, CLUSTERS, L, FUNC, ALPHA_VAL, SVD_METHOD)
computes the rank-L approximation X*Y of the input matrix
A with the Clustered Latent Semantic Indexing Method [1],
using the cluster structure information from CLUSTERS.
FUNC denotes the method used for the selection of the
number of factors from each cluster. Possible values for
FUNC:

- 'f': Selection using a heuristic method from [1]
  (see KS_SELECTION).
- 'f1': Same as 'f' but use at least one factor
  from each cluster.
- 'equal': Use the same number of factors from
  each cluster.

ALPHA_VAL is a value in [0, 1] used in the number of
factors selection heuristic [1]. Finally, SVD_METHOD
defines the method used for the computation of the SVD
(svds or propack).

REFERENCES:

[1] D. Zeimpekis and E. Gallopoulos. CLSI: A Flexible
Approximation Scheme from Clustered Term-Document Matrices.
In Proc. 5th SIAM International Conference on Data Mining,
pages 631635, Newport Beach, California, 2005.

clustering_gui
CLUSTERING_GUI CLUSTERING_GUI is a graphical user interface for all clustering functions of the Text to Matrix Generator (TMG) Toolbox.

cm

CM - computes a rank-L approximation of the input matrix
using the Centroids Method [1]

[X, Y]=CM(A, CLUSTERS) computes the rank-K approximation
X*Y of the input matrix A with the Centroids Method [1],
using the cluster structure information from CLUSTERS.

REFERENCES:
[1] H. Park, M. Jeon, and J. Rosen. Lower Dimensional Representation
of Text Data Based on Centroids and Least Squares.
BIT Numerical Mathematics, 43(2):427448, 2003.

col_normalization
COL_NORMALIZATION - normalizes the columns of the input matrix.

col_rearrange
COL_REARRANGE - reorders a matrix using a clustering result [A, N_COLS, COL_INDS]=COL_REARRANGE(A, CLUSTERS) reorders the columns of matrix A using the clustering result represented by the structure CLUSTERS. N_COLS stores the last column index for each column block, while COL_INDS containes the permuted column indices.

comppute_scat
COMPUTE_SCAT - computes the cluster selection criterion value of PDDP SCAT=COMPUTE_SCAT(A, C) returns the square of the frobenius norm of A-C*ones(1, size(A, 2)).

create_kmeans_response
CREATE_KMEANS_RESPONSE returns an html response for k-means CREATE_KMEANS_RESPONSE(CLUSTERS, TITLES) creates a summary html file containing information for the result of the k-means algorithm, defined by CLUSTERS, when applied to the dataset with document titles defined in the TITLES cell array. CREATE_KMEANS_RESPONSE(CLUSTERS, TITLES, VARIANT) defines additionaly the k-means variant (possible values 'k-means' and 'skmeans'). The result is stored in the "results" directory and displayed using the default web browser.

`create_pddp_response`

CREATE_PDDP_RESPONSE returns an html response for PDDP
CREATE_PDDP_RESPONSE(TREE_STRUCT, CLUSTERS, L, TITLES)
creates a summary html file containing information for
the result of the PDDP algorithm, defined by TREE_STRUCT
and CLUSTERS, when applied to the dataset with document
titles defined in the TITLES cell array. L defines the
maximum number of principal directions used by PDDP.
The result is stored in the ”results” directory and
displayed using the default web browser.

create_retrieval_response
CREATE_RETRIEVAL_RESPONSE returns an html response for a query CREATE_RETRIEVAL_RESPONSE(DATASET, IDS, SIMILARITY, QUERY) creates an html file containing information for the text of documents of DATASET stored in MySQL defined by IDS and having SIMILARITY similarity coefficients against QUERY. The result is stored in the ”results” directory and displayed using the default web browser.

diff_vector
DIFF_VECTOR DIFF_VECTOR returns the vector of differences between consecutive elements of the input vector.

dr_gui
DR_GUI DR_GUI is a graphical user interface for all dimensionality reduction functions of the Text to Matrix Generator (TMG) Toolbox.

ekmeans

EKMEANS - Euclidean k-Means Clustering Algorithm

EKMEANS clusters a term-document matrix using the standard k-means clustering algorithm. CLUSTERS=EKMEANS(A, C, K, TERMINATION) returns a cluster structure with K clusters for the term-document matrix A using as initial centroids the columns of C (initialized randomly when it is empty). TERMINATION defines the termination method used in k-means ('epsilon' stops iteration when objective function decrease falls down a user defined threshold - see OPTIONS input argument - while 'n_iter' stops iteration when a user defined number of iterations has been reached). [CLUSTERS, Q]=EKMEANS(A, C, K, TERMINATION) returns also the vector of objective function values for each iteration and [CLUSTERS, Q, C]=EKMEANS(A, C, K, TERMINATION) returns the final centroid vectors. EKMEANS(A, C, K, TERMINATION, OPTIONS) defines optional parameters:

- OPTIONS.iter: Number of iterations (default 10).
- OPTIONS.epsilon: Value for epsilon convergence criterion (default 1).
- OPTIONS.dsp: Displays results (default 1) or not (0) to the command window.

entropy
ENTROPY - computes the entropy of a clustering result [VENTROPY, CONFUSION_MATRIX, MISTAKES]=ENTROPY(CLUSTERS, LABELS) computes the entropy value of a clustering result represented by the CLUSTERS structure. LABELS is a vector of integers containing the true labeling of the objects. The entropy value is stored in VENTOPY, while CONFUSION_MATRIX is a k x r matrix, where k is the number of clusters and r the number of true classes, and CONFUSION_MATRIX(i, j) records the number of objects of class j assigned to cluster i. Finally, MISTAKES contains the number of misassigned objects, measured by m1+...+mk, where mi=sum(CONFUSION_MATRIX(i, j)), j∼=i.

get_node_scat
GET_NODE_SCAT - returns the PDDP node with the maximum scatter value (see PDDP) [MAX_SCAT_IND, M_SCAT]=GET_NODE_SCAT(TREE_STRUCT, SPLITTED) returns the node index and the scatter value of the PDDP tree defined by TREE_STRUCT. SPLITTED is a vector that determines the active nodes.

gui
GUI GUI is a simple, top graphical user interface of the Text to Matrix Generator (TMG) Toolbox. Using GUI, the user can select any of the four GUI modules (indexing, dimensionality reduction, clustering, classification) of TMG.

init_tmg
INIT_TMG - Installation script of TMG INIT_TMG is the installation script of the Text to Matrix Generator (TMG) Toolbox. INIT_TMG creates the MySQL database and adds all TMG directories to the path.

knn_multi
KNN_MULTI - k-Nearest Neighboors classifier for multi-label collections LABELS_AS=KNN_MULTI(A, Q, K, LABELS, NORMALIZED_DOCS, THRESHOLDS) classifies the columns of Q with the K-Nearest Neighboors classifier using the pre-classified columns of matrix A with labels LABELS (cell array of vectors of integers). THRESHOLDS is a vector of class threshold values. NORMALIZED_DOCS defines if cosine (1) or euclidean distance (0) similarity measure is to be used. LABELS_AS contains the assigned labels for the columns of Q.

knn_single
KNN_SINGLE - k-Nearest Neighboors classifier for single-label collections LABELS_AS=KNN_SINGLE(A, Q, K, LABELS, NORMALIZED_DOCS) classifies the columns of Q with the K-Nearest Neighboors classifier using the pre-classified columns of matrix A with labels LABELS (vector of integers). NORMALIZED_DOCS defines if cosine (1) or euclidean distance (0) similarity measure is to be used. LABELS_AS contains the assigned labels for the columns of Q.

ks_selection
KS_SELECTION - implements the heuristic method from [2] for the selection of the number of factors from each cluster used in the Clustered Latent Semantic Indexing method [1]. N_ST=KS_SELECTION(A, N_COLS, ALPHA_VAL, L) returns in N_ST a vector of integers denoting the number of factors (sum equals L) selected from each cluster of the tdm A. N_COLS is a vector containing the last column index for each column block, while ALPHA_VAL is a value in [0, 1].

ks_selection1
KS_SELECTION1 - implements the heuristic method from [2] for the selection of the number of factors from each cluster used in the Clustered Latent Semantic Indexing method [1]. The number of factors from each cluster is at least 1. N_ST=KS_SELECTION1(A, N_COLS, ALPHA_VAL, L) returns in N_ST a vector of integers denoting the number of factors (sum equals L) selected from each cluster of the tdm A. N_COLS is a vector containing the last column index for each column block, while ALPHA_VAL is a value in [0, 1].

`llsf_multi`

LLSF_MULTI - Linear Least Squares Fit for multi-label
collections [2]

LABELS_AS=LLSF_MULTI(A, Q, CLUSTERS, LABELS, L, METHOD,
THRESHOLDS, SVD_METHOD, CLSI_METHOD) classifies the
columns of Q with the Linear Least Squares Fit classifier
[2] using the pre-classified columns of matrix A with
labels LABELS (cell array of vectors of integers).
THRESHOLDS is a vector of class threshold values, while
CLUSTERS is a structure defining the classes. METHOD
is the method used for the approximation of the rank-l
truncated SVD, with possible values:

- 'clsi': Clustered Latent Semantic Indexing [3].
- 'cm': Centroids Method [1].
- 'svd': Singular Value Decomosition.

SVD_METHOD defines the method used for the computation
of the SVD, while CLSI_METHOD defines the method used
for the determination of the number of factors from each
class used in Clustered Latent Semantic Indexing in case
METHOD equals 'clsi'.

llsf_single

LLSF_SINGLE - Linear Least Squares Fit for single-label
collections [2]

LABELS_AS=LLSF_SINGLE(A, Q, CLUSTERS, LABELS, L, METHOD,
SVD_METHOD, CLSI_METHOD) classifies the columns of Q
with the Linear Least Squares Fit classifier [2] using
the pre-classified columns of matrix A with labels LABELS
(cell array of vectors of integers). CLUSTERS is a
structure defining the classes. METHOD is the method
used for the approximation of the rank-l truncated SVD,
with possible values:

- 'clsi': Clustered Latent Semantic Indexing [3].
- 'cm': Centroids Method [1].
- 'svd': Singular Value Decomosition.

SVD_METHOD defines the method used for the computation
of the SVD, while CLSI_METHOD defines the method used for
the determination of the number of factors from each class
used in Clustered Latent Semantic Indexing in case METHOD
equals 'clsi'.

lsa
LSA - Applies the Latent Semantic Analysis Model to a document collection [SC, DOCS_INDS] = LSA(D, P, Q, NORMALIZE_DOCS) applies LSA to the text collection represented by the latent semantic factors D, P of the collection's term - document matrix, for the query defined by the vector Q [1]. NORMALIZE_DOCS defines if the document vectors are to be normalized (1) or not (0). SC contains the sorted similarity coefficients, while DOC_INDS contains the corresponding document indices.

make_clusters_multi
MAKE_CLUSTERS_MULTI - auxiliary function for the classification algorithms CLUSTERS=MAKE_CLUSTERS_MULTI(LABELS) forms the cluster structure of a multi-label collection with document classes defined by LABELS (cell array of vectors of integers).

make_clusters_single

MAKE_CLUSTERS_SINGLE - auxiliary function for the classification algorithms

CLUSTERS=MAKE_CLUSTERS_SINGLE(LABELS) forms the cluster structure of a single-label collection with document classes defined by LABELS (vector of integers).

make_labels

MAKE_LABELS - creates a label vector of integers for the input
cell array of string

[LABELS, UNIQUE_LABELS]=MAKE_LABELS(INPUT_LABELS) creates a vector of integer labels (LABELS) for the input cell array of strings INPUT_LABELS. UNIQUE_LABELS contains the strings of unique labels of the input cell array.

make_val_inds

MAKE_VAL_INDS - auxiliary function for the classification algorithms

INDS=MAKE_VAL_INDS(LABELS) constructs an index vector used during the thresholding phase of any classifier for the multi-label collection with document classes defined by LABELS (cell array of vectors of integers).

merge_dictionary

MERGE_DICTIONARY - merges two cell arrays of chars and returns
only the distinct elements of their union (used by tmg.m,
tmg_query.m, tdm_update.m)
[ALL_WORDS, ALL_DOC_IDS]=MERGE_DICTIONARY(ALL_WORDS,
NEW_WORDS,
ALL_DOC_IDS, NEW_DOC_IDS) returns in ALL_WORDS all distinct
elements of the union of the cell arrays of chars ALL_WORDS,
NEW_WORDS corresponding to two document collections. ALL_DOC_IDS
and NEW_DOC_IDS contain the inverted indices of the two
collections. Output argument ALL_DOC_IDS contains the inverted
index of the whole collection.

`merge_tdms`

```
MERGE_TDMS - Merges two document collections
    [A, DICTIONARY]=MERGE_TDMS(A1, DICTIONARY1, A2, DICTIONARY2]
    merges the tdms A1 and A2 with corresponding dictionaries
    DICTIONARY1 and DICTIONARY2.
    MERGE_TDS(A1, DICTIONARY1, A2, DICTIONARY2, OPTIONS) defines
    optional parameters:
        - OPTIONS.min_local_freq: The minimum local frequency for
          a term (default 1)
        - OPTIONS.max_local_freq: The maximum local frequency for
          a term (default inf)
        - OPTIONS.min_global_freq: The minimum global frequency
          for a term (default 1)
        - OPTIONS.max_global_freq: The maximum global frequency
          for a term (default inf)
        - OPTIONS.local_weight: The local term weighting function
          (default 't'). Possible values (see [1, 2]):
              't': Term Frequency
              'b': Binary
              'l': Logarithmic
              'a': Alternate Log
              'n': Augmented Normalized Term Frequency
        - OPTIONS.global_weight: The global term weighting function
          (default 'x'). Possible values (see [1, 2]):
              'x': None
              'e': Entropy
              'f': Inverse Document Frequency (IDF)
              'g': GfIdf
              'n': Normal
              'p': Probabilistic Inverse
        - OPTIONS.normalization: Indicates if we normalize the
          document vectors (default 'x'). Possible values:
              'x': None
              'c': Cosine
```

myperms

MYPERMS - computes all possible combinations of the input
V=MYPERMS[P, L] returns all possible combinations of the
input vector of integers with L numbers.

open_file
OPEN_FILE OPEN_FILE is a graphical user interface for selecting a file, directory or variable from the workspace. The function returns the name of the selected file, directory or variable.

opt_2means
OPT_2MEANS - a special case of k-means for k=2 OPT_2MEANS(A, X) returns the clustering that optimizes the objective function of the k-means algorithm based on the ordering of vector X. [CLUSTERS, S]=OPT_2MEANS(A, X) returns the cluster structure as well as the value of the objective function.

`pca`

PCA - Principal Component Analysis

[U, S, V]=PCA(A, C, K, METHOD) computes the K-factor
Principal Component Analysis of A, i.e. SVD of
A-C*ones(size(A, 2), 1), using either the svds function
of MATLAB or the PROPACK package [1].

REFERENCES:
[1] R.M.Larsen, PROPACK: A Software Package for the Symmetric Eigenvalue
Problem and Singular Value Problems on Lanczos and Lanczos Bidiagonalization
with Partial Reorthogonalization, Stanford University,
http://sun.stanford.edu/∼rmunk/PROPACK.

pca_mat
PCA_MAT - Principal Component Analysis with MATLAB (svds) [U, S, V]=PCA_MAT(A, C, K) computes the K-factor Principal Component Analysis of A, i.e. SVD of A-C*ones(size(A, 2), 1), using the svds function of MATLAB.

pca_mat_afun
PCA_MAT_AFUN - Auxiliary function used in PCA_MAT.

pca_propack

PCA_PROPACK - Principal Component Analysis with PROPACK

[U, S, V]=PCA_PROPACK(A, C, K) computes the K-factor Principal Component Analysis of A, i.e. SVD of A-C*ones(size(A, 2), 1), using the PROPACK package [1].

REFERENCES:

[1] R.M.Larsen, PROPACK: A Software Package for the Symmetric Eigenvalue Problem and Singular Value Problems on Lanczos and Lanczos Bidiagonalization with Partial Reorthogonalization, Stanford University, http://sun.stanford.edu/ rmunk/PROPACK.

pca_propack_Atransfunc
PCA_PROPACK_ATRANSFUNC - Auxiliary function used in PCA_PROPACK.

pca_propack_afun
PCA_PROPACK_AFUN - Auxiliary function used in TMG_PCA_PROPACK.

pddp

PDDP - Principal Direction Divisive Partitioning Clustering Algorithm

PDDP clusters a term-document matrix (tdm) using the Principal Direction Divisive Partitioning clustering algorithm [1, 2].
CLUSTERS=PDDP(A, K, L) returns a cluster structure with K clusters for the tdm A formed using information from the first L principal components of the tdm.
[CLUSTERS, TREE_STRUCT]=PDDP(A, K, L) returns also the full PDDP tree, while [CLUSTERS, TREE_STRUCT, S]=PDDP(A, K, L) returns the objective function of PDDP.
PDDP(A, K, L, SVD_METHOD) defines the method used for the computation of the PCA (svds - default - or propack), while PDDP(A, K, L, SVD_METHOD, DSP) defines if results are to be displayed to the command window (default 1) or not (0).

REFERENCES:
[1] D.Boley, Principal Direction Divisive Partitioning, Data Mining and Knowledge Discovery 2 (1998), no. 4, 325-344.
[2] D.Zeimpekis, E.Gallopoulos, PDDP(l): Towards a Flexible Principal Direction Divisive Partitioning Clustering Algorithmm, Proc. IEEE ICDM'03 Workshop on Clustering Large Data Sets (Melbourne, Florida), 2003.

pddp_2means

PDDP_2MEANS - Hybrid Principal Direction Divisive Partitioning Clustering Algorithm and k-means

PDDP_2MEANS clusters a term-document matrix (tdm) using a combination of the Principal Direction Divisive Partitioning clustering algorithm [1] and k-means [2].
CLUSTERS=PDDP_2MEANS(A, K) returns a cluster structure with K clusters for the tdm A.
[CLUSTERS, TREE_STRUCT]=PDDP_2MEANS(A, K) returns also the full PDDP tree, while [CLUSTERS, TREE_STRUCT, S]=PDDP_2MEANS(A, K) returns the objective function of PDDP.
PDDP_2MEANS(A, K, SVD_METHOD) defines the method used for the computation of the PCA (svds - default - or propack).
PDDP_2MEANS(A, K, SVD_METHOD, DSP) defines if results are to be displayed to the command window (default 1) or not (0).
Finally, PDDP_2MEANS(A, K, SVD_METHOD, DSP, EPSILON)defines the termination criterion value for the k-means algorithm.

REFERENCES:
[1] D.Boley, Principal Direction Divisive Partitioning, Data Mining and Knowledge Discovery 2 (1998), no. 4, 325-344.
[2] D.Zeimpekis, E.Gallopoulos, k-means Steering of Spectral Divisive Clustering Algorithms, Proc. of Text Mining Workshop, Minneapolis, 2007.

pddp_optcut

PDDP_OPTCUT - Hybrid Principal Direction Divisive
Partitioning Clustering Algorithm and k-means

PDDP_OPTCUT clusters a term-document matrix (tdm) using
a combination of the Principal Direction Divisive
Partitioning clustering algorithm [1] and k-means [2].
CLUSTERS=PDDP_OPTCUT(A, K) returns a cluster structure
with K clusters for the tdm A.
[CLUSTERS, TREE_STRUCT]=PDDP_OPTCUT(A, K) returns also the
full PDDP tree, while [CLUSTERS, TREE_STRUCT, S]=PDDP_OPTCUT(A,
K) returns the objective function of PDDP.
PDDP_OPTCUT(A, K, SVD_METHOD) defines the method used for the
computation of the PCA (svds - default - or propack).
PDDP_OPTCUT(A, K, SVD_METHOD, DSP) defines if results are to be
displayed to the command window (default 1) or not (0). Finally,
PDDP_OPTCUT(A, K, SVD_METHOD, DSP, EPSILON) defines the
termination criterion value for the k-means algorithm.

REFERENCES:
[1] D.Boley, Principal Direction Divisive Partitioning, Data
Mining and Knowledge Discovery 2 (1998), no. 4, 325-344.
[2] D.Zeimpekis, E.Gallopoulos, k-means Steering of Spectral
Divisive Clustering Algorithms, Proc. of Text Mining Workshop,
Minneapolis, 2007.

`pddp_optcut_2means`

PDDP_OPTCUT_2MEANS - Hybrid Principal Direction Divisive Partitioning Clustering Algorithm and k-means

PDDP_OPTCUT_2MEANS clusters a term-document matrix (tdm) using a combination of the Principal Direction Divisive Partitioning clustering algorithm [1] and k-means [2].
CLUSTERS=PDDP_OPTCUT_OPTCUT_2MEANS(A, K) returns a cluster structure with K clusters for the tdm A.
[CLUSTERS, TREE_STRUCT]=PDDP_OPTCUT_2MEANS(A, K) returns also the full PDDP tree, while [CLUSTERS, TREE_STRUCT, S]= PDDP_OPTCUT_2MEANS(A, K) returns the objective function of PDDP.
PDDP_OPTCUT_2MEANS(A, K, SVD_METHOD) defines the method used for the computation of the PCA (svds - default - or propack).
PDDP_OPTCUT_2MEANS(A, K, SVD_METHOD, DSP) defines if results are to be displayed to the command window (default 1) or not (0). Finally, PDDP_OPTCUT_2MEANS(A, K, SVD_METHOD, DSP, EPSILON) defines the termination criterion value for the k-means algorithm.

REFERENCES:
[1] D.Boley, Principal Direction Divisive Partitioning, Data Mining and Knowledge Discovery 2 (1998), no. 4, 325-344.
[2] D.Zeimpekis, E.Gallopoulos, k-means Steering of Spectral Divisive Clustering Algorithms, Proc. of Text Mining Workshop, Minneapolis, 2007.

pddp_optcutpd

PDDP_OPTCUTPD - Hybrid Principal Direction Divisive Partitioning Clustering Algorithm and k-means

PDDP_OPTCUTPD clusters a term-document matrix (tdm) using a combination of the Principal Direction Divisive Partitioning clustering algorithm [1, 2] and k-means [3].
CLUSTERS=PDDP_OPTCUT_OPTCUTPD(A, K, L) returns a cluster structure with K clusters for the tdm A formed using information from the first L principal components of the tdm.
[CLUSTERS, TREE_STRUCT]=PDDP_OPTCUTPD(A, K, L) returns also the full PDDP tree, while [CLUSTERS, TREE_STRUCT, S]= PDDP_OPTCUTPD(A, K, L) returns the objective function of PDDP. PDDP_OPTCUTPD(A, K, L, SVD_METHOD) defines the method used for the computation of the PCA (svds - default - or propack). Finally, PDDP_OPTCUTPD(A, K, L, SVD_METHOD, DSP) defines if results are to be displayed to the command window (default 1) or not (0).

REFERENCES:
[1] D.Boley, Principal Direction Divisive Partitioning, Data Mining and Knowledge Discovery 2 (1998), no. 4, 325-344.
[2] D.Zeimpekis, E.Gallopoulos, PDDP(l): Towards a Flexible Principal Direction Divisive Partitioning Clustering Algorithmm, Proc. IEEE ICDM'03 Workshop on Clustering Large Data Sets (Melbourne, Florida), 2003.
[3] D.Zeimpekis, E.Gallopoulos, k-means Steering of Spectral Divisive Clustering Algorithms, Proc. of Text Mining Workshop, Minneapolis, 2007.

ps_pdf2ascii
PS_PDF2ASCII - converts the input ps or pdf file to ASCII RESULT = PS_PDF2ASCII(FILENAME) converts the input ps or pdf files to ASCII, using ghostscript's utility 'ps2ascii'. RESULT returns a success indicator, e.g. -2 if the input file does not exist or has a wrong format, -1 if gs is not installed or the path isn't set, 0 if 'ps2ascii' didn't work properly, and 1 if the conversion was successful.

retrieval_gui
RETRIEVAL_GUI RETRIEVAL_GUI is a graphical user interface for all retrieval functions of the Text to Matrix Generator (TMG) Toolbox.

rocchio_multi
ROCCHIO_MULTI - Rocchio classifier for multi-label collections LABELS_AS=KNN_MULTI(A, CLUSTERS, BETA, GAMMA, Q, LABELS, NORMALIZED_DOCS, THRESHOLDS) classifies the columns of Q with the Rocchio classifier using the pre-classified columns of matrix A with labels LABELS (vector of integers). THRESHOLDS is a vector of class threshold values. BETA and GAMMA define the weight of positive and negative examples in the formation of each class centroid. NORMALIZED_DOCS defines if cosine (1) or euclidean distance (0) similarity measure is to be used. LABELS_AS contains the assigned labels for the columns of Q.

rocchio_single
ROCCHIO_SINGLE - Rocchio classifier for single-label collections LABELS_AS=KNN_SINGLE(A, CLUSTERS, BETA, GAMMA, Q, LABELS, NORMALIZED_DOCS) classifies the columns of Q with the Rocchio classifier using the pre-classified columns of matrix A with labels LABELS (vector of integers). BETA and GAMMA define the weight of positive and negative examples in the formation of each class centroid. NORMALIZED_DOCS defines if cosine (1) or euclidean distance (0) similarity measure is to be used. LABELS_AS contains the assigned labels for the columns of Q.

scut_knn

SCUT_KNN - implements the Scut thresholding technique from [1]
for the k-Nearest Neighboors classifier

THRESHOLD=SCUT_KNN(A, Q, K, LABELS_TR, LABELS_TE, MINF1,
NORMALIZE, STEPS) returns the vector of thresholds for
the k-Nearest Neighboors classifier for the collection
[A Q]. A and Q define the training and test parts of
the validation set with labels LABELS_TR and LABELS_TE
respectively. MINF1 defines the minimum F1 value and
NORMALIZE defines if cosine (1) or euclidean distance (0)
measure of similarity is to be used. Finally, STEPS
defines the number of steps used during thresholding.
[THRESHOLD, F, THRESHOLDS]=SCUT_KNN(A, Q, K, LABELS_TR,
LABELS_TE, MINF1, NORMALIZE, STEPS) returns also the best
F1 value as well as the matrix of thresholds for each step
(row i corresponds to step i).

REFERENCES:
[1] Y. Yang. A Study of Thresholding Strategies for Text
Categorization. In Proc. 24th ACM SIGIR, pages 137145,
New York, NY, USA, 2001. ACM Press.

scut_llsf

SCUT_LLSF - implements the Scut thresholding technique from [2]
for the Linear Least Squares Fit classifier [3]

THRESHOLD=SCUT_LLSF(A, Q, CLUSTERS, K, LABELS_TR, LABELS_TE, MINF1, L, METHOD, STEPS, SVD_METHOD, CLSI_METHOD) returns the vector of thresholds for the Linear Least Squares Fit classifier for the collection [A Q]. A and Q define the training and test parts of the validation set with labels LABELS_TR and LABELS_TE respectively. CLUSTERS is a structure defining the classes, while MINF1 defines the minimum F1 value and STEPS defines the number of steps used during thresholding.
METHOD is the method used for the approximation of the rank-l truncated SVD, with possible values:

- 'clsi': Clustered Latent Semantic Indexing [4].
- 'cm': Centroids Method [1].
- 'svd': Singular Value Decomosition.

SVD_METHOD defines the method used for the computation of the SVD, while CLSI_METHOD defines the method used for the determination of the number of factors from each class used in Clustered Latent Semantic Indexing in case METHOD equals 'clsi'.
[THRESHOLD, F, THRESHOLDS]=SCUT_LLSF(A, Q, CLUSTERS, K, LABELS_TR, LABELS_TE, MINF1, L, METHOD, STEPS, SVD_METHOD, CLSI_METHOD) returns also the best F1 value as well as the matrix of thresholds for each step (row i corresponds to step i).

REFERENCES:

[1] H. Park, M. Jeon, and J. Rosen. Lower Dimensional Representation of Text Data Based on Centroids and Least Squares. BIT Numerical Mathematics, 43(2):427448, 2003.
[2] Y. Yang. A Study of Thresholding Strategies for Text Categorization. In Proc. 24th ACM SIGIR, pages 137145, New York, NY, USA, 2001. ACM Press.
[3] Y. Yang and C. Chute. A Linear Least Squares Fit Mapping Method for Information Retrieval from Natural Language Texts. In Proc. 14th Conference on Computational Linguistics, pages 447453, Morristown, NJ, USA, 1992.
[4] D. Zeimpekis and E. Gallopoulos, "Non-Linear Dimensional Reduction via Class Representatives for Text Classification". In Proc. 2006 IEEE International Conference on Data Mining (ICDM'06), Hong Kong, Dec. 2006.

scut_rocchio

SCUT_ROCCHIO - implements the Scut thresholding technique
from [1] for the Rocchio classifier

THRESHOLD=SCUT_ROCCHIO(A, CLUSTERS, BETA, GAMMA, Q,
LABELS_TR, LABELS_TE, MINF1, NORMALIZE, STEPS) returns
the vector of thresholds for the Rocchio classifier
for the collection [A Q]. A and Q define the training
and test parts of the validation set with labels
LABELS_TR and LABELS_TE respectively. MINF1 defines
the minimum F1 value, while NORMALIZE defines if cosine
(1) or euclidean distance (0) measure of similarity is
to be used, CLUSTERS is a structure defining the classes
and STEPS defines the number of steps used during
thresholding. BETA and GAMMA define the weight of positive
and negative examples in the formation of each class
centroid.

[THRESHOLD, F, THRESHOLDS]=SCUT_ROCCHIO(A, CLUSTERS, BETA,
GAMMA, Q, LABELS_TR, LABELS_TE, MINF1, NORMALIZE, STEPS)
returns also the best F1 value as well as the matrix of
thresholds for each step (row i corresponds to step i).

REFERENCES:

[1] Y. Yang. A Study of Thresholding Strategies for Text
Categorization. In Proc. 24th ACM SIGIR, pages 137145,
New York, NY, USA, 2001. ACM Press.

sdd_tmg
SDD_TMG - interface for SDDPACK [X, D, Y]=SDD_TMG(A, K) computes a rank-K Semidiscrete Decomposition of A using the SDDPACK [1]. REFERENCES: Tamara G. Kolda and Dianne P. O'Leary, Computation and Uses of the Semidiscrete Matrix Decomposition, Computer Science Department Report CS-TR-4012 Institute for Advanced Computer Studies Report UMIACS-TR-99-22, University of Maryland, April 1999.

`skmeans`

SKMEANS - Spherical k-Means Clustering Algorithm

SKMEANS clusters a term-document matrix using the Spherical k-means clustering algorithm [1]. CLUSTERS=SKMEANS(A, C, K, TERMINATION) returns a cluster structure with K clusters for the term-document matrix A using as initial centroids the columns of C (initialized randomly when it is empty). TERMINATION defines the termination method used in spherical k-means ('epsilon' stops iteration when objective function increase falls down a user defined threshold - see OPTIONS input argument - while 'n_iter' stops iteration when a user defined number of iterations has been reached). [CLUSTERS, Q]=SKMEANS(A, C, K, TERMINATION) returns also the vector of objective function values for each iteration and [CLUSTERS, Q, C]=SKMEANS(A, C, K, TERMINATION) returns the final centroid vectors. SKMEANS(A, C, K, TERMINATION, OPTIONS) defines optional parameters:

- OPTIONS.iter: Number of iterations (default 10).
- OPTIONS.epsilon: Value for epsilon convergence criterion (default 1).
- OPTIONS.dsp: Displays results (default 1) or not (0) to the command window.

REFERENCES:

[1] I. S. Dhillon and D. M. Modha, "Concept Decompositions for Large Sparse Text Data using Clustering", Machine Learning, 42:1, pages 143-175, Jan, 2001.

stemmer

STEMMER - applies the Porter's Stemming algorithm [1]

S = STEMMER(TOKEN, DSP) returns in S the stemmed word of TOKEN. DSP indicates if the function displays the result of each stem (1).

REFERENCES:

[1] M.F.Porter, An algorithm for suffix stripping, Program, 14(3): 130-137, 1980.

strip_html
STRIP_HTML - removes html entities from an html file S = STRIP_HTML(FILENAME) parses file FILENAME and removes the html entities, while the result is stored in S as a cell array and written in file ”FILENAME.TXT”.

svd_tmg
SVD_TMG - Singular Value Decomposition [U, S, V]=SVD_TMG(A, K, METHOD) computes the K-factor truncated Singular Value Decomposition of A using either the svds function of MATLAB or the PROPACK package [1]. REFERENCES: [1] R.M.Larsen, PROPACK: A Software Package for the Symmetric Eigenvalue Problem and Singular Value Problems on Lanczos and Lanczos Bidiagonalization with Partial Reorthogonalization, Stanford University, http://sun.stanford.edu/ rmunk/PROPACK.

tdm_downdate

TDM_DOWNDATE - renews a text collection by downdating the correspoding term-document matrix

A = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns the new term - document matrix of the downdated collection. UPDATE_STRUCT defines the update structure returned by TMG, while REMOVED_DOCS defines the indices of the documents that is to be be removed.

[A, DICTIONARY] = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns also the dictionary for the updated collection, while

[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZED_FACTORS] = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns the vectors of global weights for the dictionary and the normalization factor for each document in case such a factor is used. If normalization is not used TDM_DOWNDATE returns a vector of all ones. [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC] = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns statistics for each document, i.e. the number of terms for each document.

[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC, TITLES, FILES] = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns in FILES the filenames containing the collection's documents and a cell array (TITLES) that contains a declaratory title for each document, as well as the document's first line.

Finally [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC, TITLES, FILES, UPDATE_STRUCT] = TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS) returns the update structure that keeps the essential information for the collection' s update (or downdate).

TDM_DOWNDATE(UPDATE_STRUCT, REMOVED_DOCS, OPTIONS) defines optional parameters:

- OPTIONS.dsp: Displays results (default 1) or not (0) to the command window.

tdm_update

TDM_UPDATE renews a text collection by updating the
correspoding term-document matrix.

A = TDM_UPDATE(FILENAME, UPDATE_STRUCT) returns the new
term - document matrix of the updated collection. FILENAME
defines the file (or files in case a directory is supplied)
containing the new documents, while UPDATE_STRUCT defines
the update structure returned by TMG. In case FILENAME
variable is empty, the collection is simply updated using
the options defined by UPDATE_STRUCT (for example, use
another term-weighting scheme).
[A, DICTIONARY] = TDM_UPDATE(FILENAME, UPDATE_STRUCT)
returns also the dictionary for the updated collection,
while [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZED_FACTORS]
= TDM_UPDATE(FILENAME, UPDATE_STRUCT) returns the vectors
of global weights for the dictionary and the normalization
factor for each document in case such a factor is used.
If normalization is not used TDM_UPDATE returns a vector
of all ones.
[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS,
WORDS_PER_DOC] = TDM_UPDATE(FILENAME, UPDATE_STRUCT) returns
statistics for each document, i.e. the number of terms for
each document.
[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS,
WORDS_PER_DOC, TITLES, FILES] = TDM_UPDATE(FILENAME,
UPDATE_STRUCT) returns in FILES the filenames contained in
directory (or file) FILENAME and a cell array (TITLES) that
containes a declaratory title for each document, as well as
the document's first line.
Finally [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS,
WORDS_PER_DOC, TITLES, FILES, UPDATE_STRUCT] =
TDM_UPDATE(FILENAME, UPDATE_STRUCT) returns the update
structure that keeps the essential information for the
collection' s update (or downdate).
TDM_UPDATE(FILENAME, UPDATE_STRUCT, OPTIONS) defines optional
parameters:

- OPTIONS.delimiter: The delimiter between documents within
  the same file. Possible values are 'emptyline' (default),
  'none_delimiter' (treats each file as a single document)
  or any other string.
- OPTIONS.line_delimiter: Defines if the delimiter takes a
  whole line of text (default, 1) or not.

- OPTIONS.update_step: The step used for the incremental built of the inverted index (default 10,000).
- OPTIONS.dsp: Displays results (default 1) or not (0) to the command window.

tmg

TMG - Text to Matrix Generator

TMG parses a text collection and generates the term - document matrix.
A = TMG(FILENAME) returns the term - document matrix, that corresponds to the text collection contained in files of directory (or file) FILENAME.
Each document must be separeted by a blank line (or another delimiter that is defined by OPTIONS argument) in each file.
[A, DICTIONARY] = TMG(FILENAME) returns also the dictionary for the collection, while [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZED_FACTORS] = TMG(FILENAME) returns the vectors of global weights for the dictionary and the normalization factor for each document in case such a factor is used. If normalization is not used TMG returns a vector of all ones.
[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC] = TMG(FILENAME) returns statistics for each document, i.e. the number of terms for each document.
[A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC, TITLES, FILES] = TMG(FILENAME) returns in FILES the filenames contained in directory (or file) FILENAME and a cell array (TITLES) that containes a declaratory title for each document, as well as the document's first line. Finally [A, DICTIONARY, GLOBAL_WEIGHTS, NORMALIZATION_FACTORS, WORDS_PER_DOC, TITLES, FILES, UPDATE_STRUCT] = TMG(FILENAME) returns a structure that keeps the essential information for the collection' s update (or downdate).

TMG(FILENAME, OPTIONS) defines optional parameters:

- OPTIONS.use_mysql: Indicates if results are to be stored in MySQL.
- OPTIONS.db_name: The name of the directory where the results are to be saved.
- OPTIONS.delimiter: The delimiter between documents within the same file. Possible values are 'emptyline' (default), 'none_delimiter' (treats each file as a single document) or any other string.
- OPTIONS.line_delimiter: Defines if the delimiter takes a whole line of text (default, 1) or not.
- OPTIONS.stoplist: The filename for the stoplist,

```
        i.e. a list of common words that we don't use for
        the indexing (default no stoplist used).
      - OPTIONS.stemming: Indicates if the stemming algorithm
        is used (1) or not (0 - default).
      - OPTIONS.update_step: The step used for the incremental
        built of the inverted index (default 10,000).
      - OPTIONS.min_length: The minimum length for a term
        (default 3).
      - OPTIONS.max_length: The maximum length for a term
        (default 30).
      - OPTIONS.min_local_freq: The minimum local frequency for
        a term (default 1).
      - OPTIONS.max_local_freq: The maximum local frequency for
        a term (default inf).
      - OPTIONS.min_global_freq: The minimum global frequency
        for a term (default 1).
      - OPTIONS.max_global_freq: The maximum global frequency
        for a term (default inf).
      - OPTIONS.local_weight: The local term weighting function
        (default 't'). Possible values (see [1, 2]):
            't': Term Frequency
            'b': Binary
            'l': Logarithmic
            'a': Alternate Log
            'n': Augmented Normalized Term Frequency
      - OPTIONS.global_weight: The global term weighting function
        (default 'x'). Possible values (see [1, 2]):
            'x': None
            'e': Entropy
            'f': Inverse Document Frequency (IDF)
            'g': GfIdf
            'n': Normal
            'p': Probabilistic Inverse
      - OPTIONS.normalization: Indicates if we normalize the
        document vectors (default 'x'). Possible values:
            'x': None
            'c': Cosine
      - OPTIONS.dsp: Displays results (default 1) or not (0) to
        the command window.

    REFERENCES:
    [1] M.Berry and M.Browne, Understanding Search Engines, Mathematical
    Modeling and Text Retrieval, Philadelphia, PA: Society for Industrial
    and Applied Mathematics, 1999.
    [2] T.Kolda, Limited-Memory Matrix Methods with Applications,
    Tech.Report CS-TR-3806, 1997.
```

tmg_gui

TMG_GUI

TMG_GUI is a graphical user interface for all indexing
routines of the Text to Matrix Generator (TMG) Toolbox.
For a full documentation type 'help tmg', 'help tmg_query',
'help tdm_update' or 'help tdm_downdate'.
For a full documentation of the GUI's usage, select the
help tab to the GUI.

tmg_query

```
TMG_QUERY - Text to Matrix Generator, query vector constructor
    TMG_QUERY parses a query text collection and generates the
    query vectors corresponding to the supplied dictionary.
    Q = TMG_QUERY(FILENAME, DICTIONARY) returns the query
    vectors, that corresponds to the text collection contained
    in files of directory FILENAME. DICTIONARY is the array of
    terms corresponding to a text collection.
    Each query must be separeted by a blank line (or another
    delimiter that is defined by OPTIONS argument) in each file.
    [Q, WORDS_PER_QUERY] = TMG_QUERY(FILENAME, DICTIONARY)
    returns statistics for each query, i.e. the number of terms
    for each query.
    Finally, [Q, WORDS_PER_QUERY, TITLES, FILES] =
    TMG_QUERY(FILENAME) returns in FILES the filenames contained
    in directory (or file) FILENAME and a cell array (TITLES)
    that containes a declaratory title for each query, as well
    as the query's first line.

    TMG_QUERY(FILENAME, DICTIONARY, OPTIONS) defines optional
    parameters:
        - OPTIONS.delimiter: The delimiter between queries within
          the same file. Possible values are 'emptyline' (default),
          'none_delimiter' (treats each file as a single query)
          or any other string.
        - OPTIONS.line_delimiter: Defines if the delimiter takes a
          whole line of text (default, 1) or not.
        - OPTIONS.stoplist: The filename for the stoplist, i.e. a
          list of common words that we don't use for the indexing
          (default no stoplist used).
        - OPTIONS.stemming: Indicates if the stemming algorithm is
          used (1) or not (0 - default).
        - OPTIONS.update_step: The step used for the incremental
          built of the inverted index (default 10,000).
        - OPTIONS.local_weight: The local term weighting function
          (default 't'). Possible values (see [1, 2]):
              't': Term Frequency
              'b': Binary
              'l': Logarithmic
              'a': Alternate Log
              'n': Augmented Normalized Term Frequenct
        - OPTIONS.global_weights: The vector of term global
          weights (returned by tmg).
        - OPTIONS.dsp: Displays results (default 1) or not (0).
```

REFERENCES:
[1] M.Berry and M.Browne, Understanding Search Engines, Mathematical Modeling and Text Retrieval, Philadelphia, PA: Society for Industrial and Applied Mathematics, 1999.
[2] T.Kolda, Limited-Memory Matrix Methods with Applications, Tech.Report CS-TR-3806, 1997.

tmg_save_results
TMG_SAVE_RESULTS TMG_SAVE_RESULTS is a graphical user interface used from TMG_GUI, for saving the results to a (or multiple) .mat file(s).

tmg_template
TDM_TEMPLATE - demo script This is a template script demonstrating the use of TMG, as well as the application of the resulting TDM'S in two IR tasks, quering and clustering. The quering models used is the Vector Space Model (see vsm.m) and LSI (see lsi.m), while two versions of the k-means algorithm (euclidean and spherical, see ekmeans.m and skmeans.m) cluster the resulting matrix (see pddp.m). The user can edit this code in order to change the default OPTIONS of TMG, as well as to apply other IR tasks or use his own implementations regarding these tasks.

two_means_1d
TWO_MEANS_1D - returns the clustering that optimizes the objective function of the k-means algorithm for the input vector. [CUTOFF, CLUSTERS, DISTANCE, OF, MEAN1, MEAN2]= TWO_MEANS_1D(A) returns the cutoff value of the clustering, the cluster structure, the separation distance, the value of the objective function and the two mean values.

unique_elements

UNIQUE_ELEMENTS - detects all distinct elements of a vector

[ELEMENTS, N] = UNIQUE_ELEMENTS(X) returns in ELEMENTS all distinct elements of vector X, and in N the number of times each element appears in X. A value is repeated if it appears in non-consecutive elements. For no repetitive elements sort the input vector.

unique_words

UNIQUE_WORDS - detects all distinct elements of a cell array of
chars (used by tmg.m, tmg_query.m, tdm_update.m)

[NEW_WORDS, NEW_DOC_IDS]=UNIQUE_WORDS(WORDS, DOC_IDS,
N_DOCS)

returns in NEW_WORDS all distinct elements of the cell array
of chars WORDS. DOC_IDS is the vector of the document identifiers
containing the corresponding words, while N_DOCS is the total
number of documents contained to the collection. NEW_DOC_IDS
contains the inverted index of the collection as a cell array
of 2 x N_DOCS arrays.

vsm

VSM - Applies the Vector Space Model to a document collection
[SC, DOCS_INDS] = VSM(D, Q, NORMALIZE_DOCS) applies the
Vector Space Model to the text collection represented by
the term - document matrix D for the query defined by the
vector Q [1]. NORMALIZE_DOCS defines if the document
vectors are to be normalized (1) or not (0). SC contains
the sorted similarity coefficients, while DOC_INDS contains
the corresponding document indices.

REFERENCES:
[1] M.Berry and M.Browne, Understanding Search Engines,
Mathematical Modeling and Text Retrieval, Philadelphia,
PA: Society for Industrial and Applied Mathematics, 1999.